

OS'25 Project

PART I: PREREQUISITES

Agenda

- PART 0: ROADMAP
- PART I: PREREQUISITES
- PART II: GROUP MODULES
- PART III: INDIVIDUAL MODULES
- PART IV: OVERALL TESTING & BONUSES

Agenda

- **PART I: PREREQUISITES**

- Pointers
- Lists
- System Calls
- Spin Locks



Pointers

PART I: PREREQUISITES

Memory and Pointers

- a. Why we need pointers?!
- b. Pointers vs. Variables...
 - i. Definition
 - ii. Data type & size
 - iii. Setting values
 - iv. Incrementing...
 - v. Structure and accessing its members

Pointer vs Variables: Definition

Pointers

Variables

```
char   *ptr;
```

```
char   x;
```

Pointer vs Variables: Data type & Size

Pointers

```
char *ptr;
```

Data type:

- Data type to point into it

Size of ptr (address):

- 4 Byte
- Size of address bus of CPU
(protected → 32-bit)

Variables

```
char x;
```

Data type:

- Data type of variable itself

Size of x:

- 1 Byte → sizeof(char)

Pointer vs Variables: Assigning value

Pointers

Variables

Assigning value:

```
char *c_ptr = 0x100; Addr.  
//changes the pointed address  
c_ptr = 0x50;  
c_ptr = &x;  
  
//changes the value within  
the pointed address  
*c_ptr = 'A';
```

Assigning value:

```
char x = 10;  
x = 50;  
x = 'A';
```

Pointer vs Variables: Incrementing

Pointers

increment:

```
char *c_ptr = 0x100;  
c_ptr++; // ptr=0x101
```

```
int *i_ptr = 0x100;  
i_ptr++; // ptr=0x104
```

Increases by the size of its type

Variables

increment:

```
char x = 10;  
x++; // x=11
```

```
int x = 10;  
x++; // x=11
```

Increases by 1

Pointer vs Variables: Structure & its members

```
struct MyStruct {  
    int x, y;  
    char c;  
    char* c_ptr;  
}
```

Pointers

Variables

Init. & Assign.:

```
Struct MyStruct *my_struct;  
my_struct->x = 5;  
my_struct->c = 'A';  
(*my_struct).y = 10
```

Init. & Assign.:

```
Struct MyStruct my_struct;  
my_struct.x = 5;  
my_struct.c = 'A';
```

Pointer vs Variables: Structure & its members



```
struct MyStruct {  
    int x, y; 8B  
    char c; 1B  
    char* c_ptr; 4B  
};
```

Pointers

Variables

Init. & Assign.:

```
Struct MyStruct *my_struct;  
my_struct = 100;  
my_struct++; // ptr=113  
increases by size of struct
```

Init. & Assign.:

```
Struct MyStruct my_struct;
```



LISTs in FOS

PART I: PREREQUISITES

How to define LISTS in FOS?

To define a LIST that points to objects of type **struct my_struct**:

1. Create a list **head** that holds info about the list (size, head, tail).

- o **LIST_HEAD**([LIST_TYPE_DEF], [STRUCT NAME THAT WILL POINTS TO]);
- o Ex: 

```
LIST_HEAD(MY_LIST_TYPE, my_struct);
```

2. Add **next** and **previous** pointers to the **struct**

- o **LIST_ENTRY**([STRUCT NAME]) prev_next_info;
- o Ex:

```
struct my_struct
{
    int x, y;
    LIST_ENTRY(my_struct) prev_next_info;
};
```

3. Define your list

- o **struct** [LIST_TYPE_DEF] my_list
- o Ex:

```
struct MY_LIST_TYPE my_list;
```

How to use LISTS?

Set of helper ready made functions are available in [Appendix](#)

- LIST_INIT (...)
- LIST_INSERT_HEAD (...)
- LIST_SIZE (...)
- LIST_REMOVE (...)
- ...

IMPORTANT: you should **pass** the list to any of these functions by **reference**
(i.e. Put **&** before the name of the list)



System Calls

PART I: PREREQUISITES

System Calls – Idea

It's OS procedure that executes privileged instructions (e.g., I/O); (API exported by kernel)

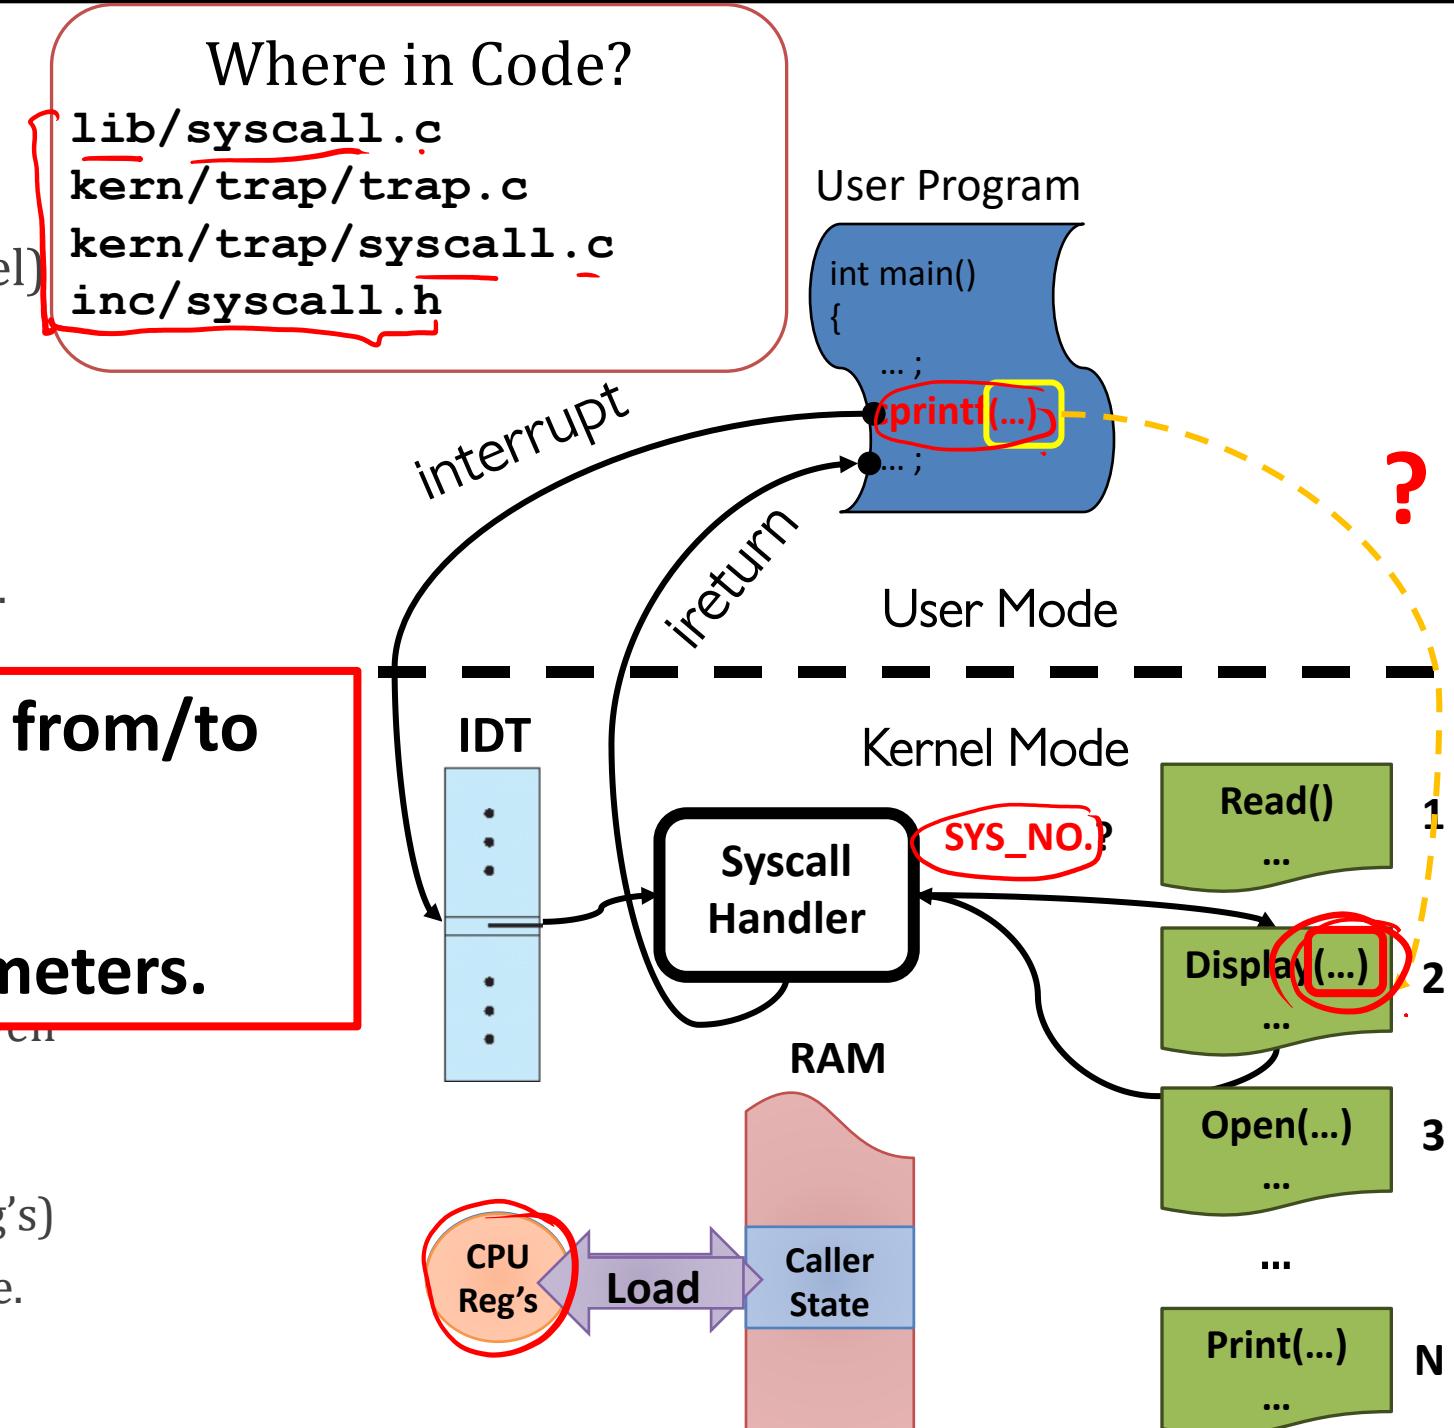
Causes a trap, which

1. Switch to the kernel mode
2. Look in Interrupt Descriptor Table (IDT)
3. Jumps to the **syscall handler** in the kernel.

How to pass params/return value from/to user to kernel or vice versa?

OS should verify the caller's parameters.

3. Call associated function that serve the given system call and pass to it the caller's **parameters**
4. After finish, **restore** caller's state (CPU Reg's)
5. use **iret** instruction to **return** to user mode.



System Calls – Params Validation

At **kernel side**: need to validate any **address (or range)** that is passed from user to kernel to ensure that:

- 1. **NOT null** pointers,
- 2. **NOT illegal** pointers (e.g. pointing to kernel memory) (i.e. outside User Area)
- 3. **NOT invalid** pointers (e.g. pointing to unmapped memory),

First **TWO cases** should be checked in the **system call**. If violated, kernel can **terminate** the **user program** by calling: **env_exit()**;

Third case cause a “page fault” that can be handled in the **page_fault()** handler.

This technique is normally faster because it takes advantage of the processor's **MMU**, so it tends to be used in real kernels (including Linux).



Spin Locks

PART I: PREREQUISITES

Locks: Definition

Locks provide two **atomic** operations:

- **Lock.acquire()** – wait until lock is **free**; then **mark** it as **busy**
 - After this returns, we say the calling thread **holds** the lock
- **Lock.release()** – **mark** lock as **free**
 - Should only be called by a thread that currently holds the lock
 - After this returns, the calling thread no longer holds the lock

Negatives of interrupt-based implementation:

- Can't give lock implementation to **users**
- Doesn't work well on **multiprocessor**
- **lock entire Building!!**

Locks: Implementation

Exchange Instruction

- Exchanges the contents of a register with a memory location.
- Available in Intel IA-32 (Pentium) and IA-64 (Itanium)
- The entire function is carried out **atomically**;
- not subject to interruption (i.e. indivisible). *By H/W*

(Swap)

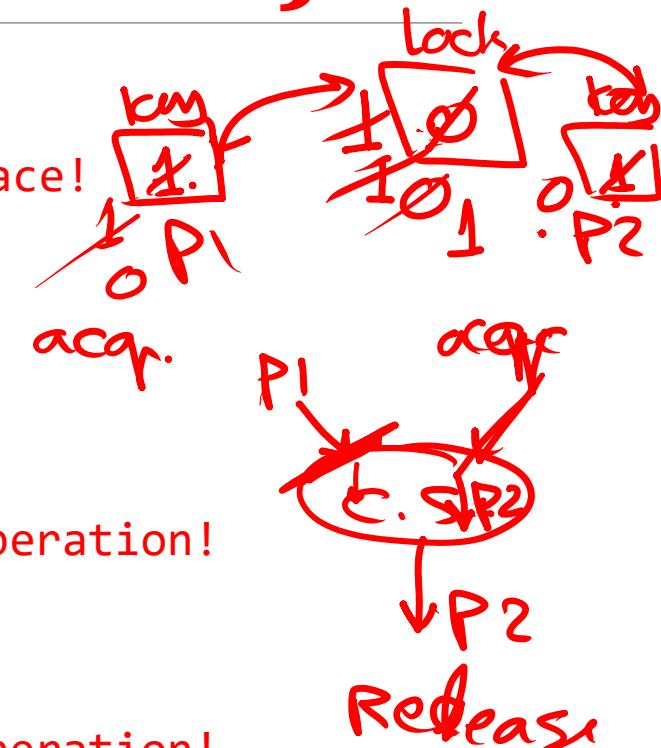
```
void exchange (int register, int memory)
{
    int temp;
    { temp = memory;
        memory = register;
        register = temp;
    }
}
```

Locks: Implementation (**SpinLock**)

Simple lock that doesn't require entry into the kernel:

```
// (Free) Can access this memory location from user space! ✓  
int lock = 0; // Interface: acquire(&lock);  
// release(&lock);  
  
acquire(int *thelock) {  
    int mykey = 1;  
    while (xchg(thelock, mykey) != 0);  
}  
  
release(int *thelock) {  
    *thelock = 0;  
} // Atomic operation!  
  
spin
```

Busy-Waiting: thread consumes cycles while waiting



Locks: Implementation (SpinLock)

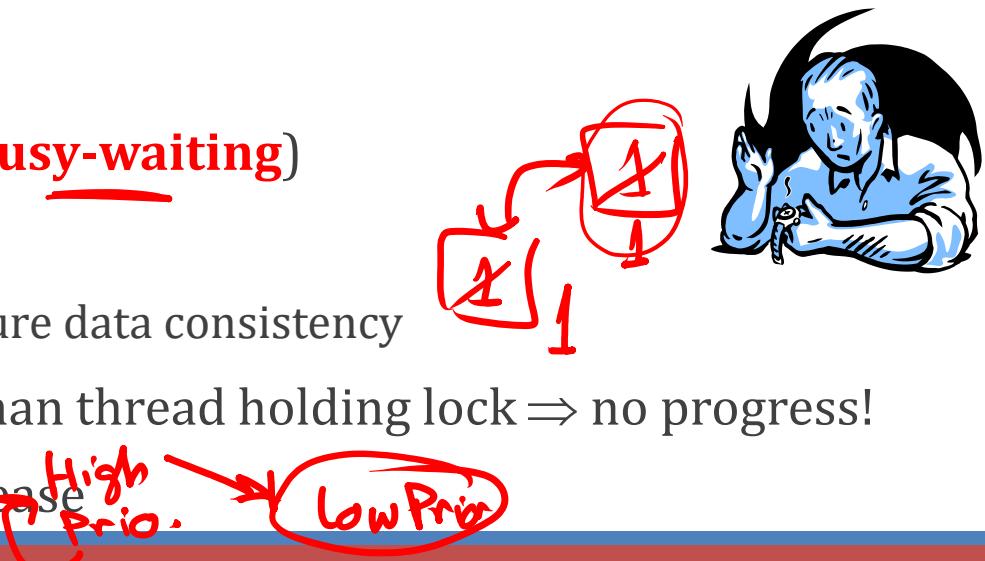
Positives

- Machine can receive interrupts
- User code can use this lock
- Works on a multiprocessor
- No System Calls at all

Usually used for
short-time critical section

Negatives

- This is very **inefficient** as thread will consume cycles waiting (**busy-waiting**)
- **Cache coherence** issue in multi-cores
 - always read-write while waiting → high traffic on cache buses to ensure data consistency
- **Priority Inversion:** If busy-waiting thread has higher priority than thread holding lock ⇒ no progress!
- **Solution:** disable the interrupt in acquire and enable it in release



SpinLock: In Code (KERNEL)

```
struct kspinlock {
    uint32 locked;           // Is the lock held?
```

```
// Acquire the lock.
// Loops (spins) until the lock is acquired.
// Holding a lock for a long time may cause
// other CPUs to waste time spinning to acquire it.
void acquire_kspinlock(struct kspinlock *lk)
{
    if(holding_kspinlock(lk))
        panic("acquire_spinlock: lock \"%s\" is already
               held");
    pushcli(); /*disable interrupts to avoid deadlock (
    // The xchg is atomic.
    while(xchg(&lk->locked, 1) != 0);
```

```
// Release the lock.
void release_kspinlock(struct kspinlock *lk)
{
    if(!holding_kspinlock(lk))
    {
        printcallstack(lk);
        panic("release: lock \"%s\" is either not held or
              already released");
    }
    // Release the lock, equivalent to lk->locked = 0.
    // This code can't use a C assignment, since it might
    // not be atomic. A real OS would use C atomics here.
    asm volatile("movl $0, %0" : "+m" (lk->locked) : );
    popcli(); //enable the interrupt
}
```

SpinLock: In Code (USER)

```
struct _uspinlock {
    uint32 locked;           // Is the lock held?
```

```
// Acquire the lock.
// Loops (spins) until the lock is acquired.
// Holding a lock for a long time may cause
// other CPUs to waste time spinning to acquire it.
void acquire_uspinlock(struct uspinlock *lk)
{
    // The xchg is atomic.
    while(xchg(&lk->locked, 1) != 0) ;
}

// Release the lock.
void release_uspinlock(struct uspinlock *lk)
{
    if(!(lk->locked))
    {
        panic("release: lock \\"%s\\" is not held!", lk->name);
    }
    // Release the lock, equivalent to lk->locked = 0.
    // This code can't use a C assignment, since it might
    // not be atomic. A real OS would use C atomics here.
    asm volatile("movl $0, %0" : "+m" (lk->locked) : );
}
```



NO INTERRUPT ACCESS @USER SIDE

NOW, LET'S START CODING TOGETHER...

😊 Enjoy developing your **own OS** 😊

