

OS'25 Project

PART II: GROUP MODULES

(DYNAMIC ALLOCATOR, KERNEL HEAP, FAULT HANDLER I)

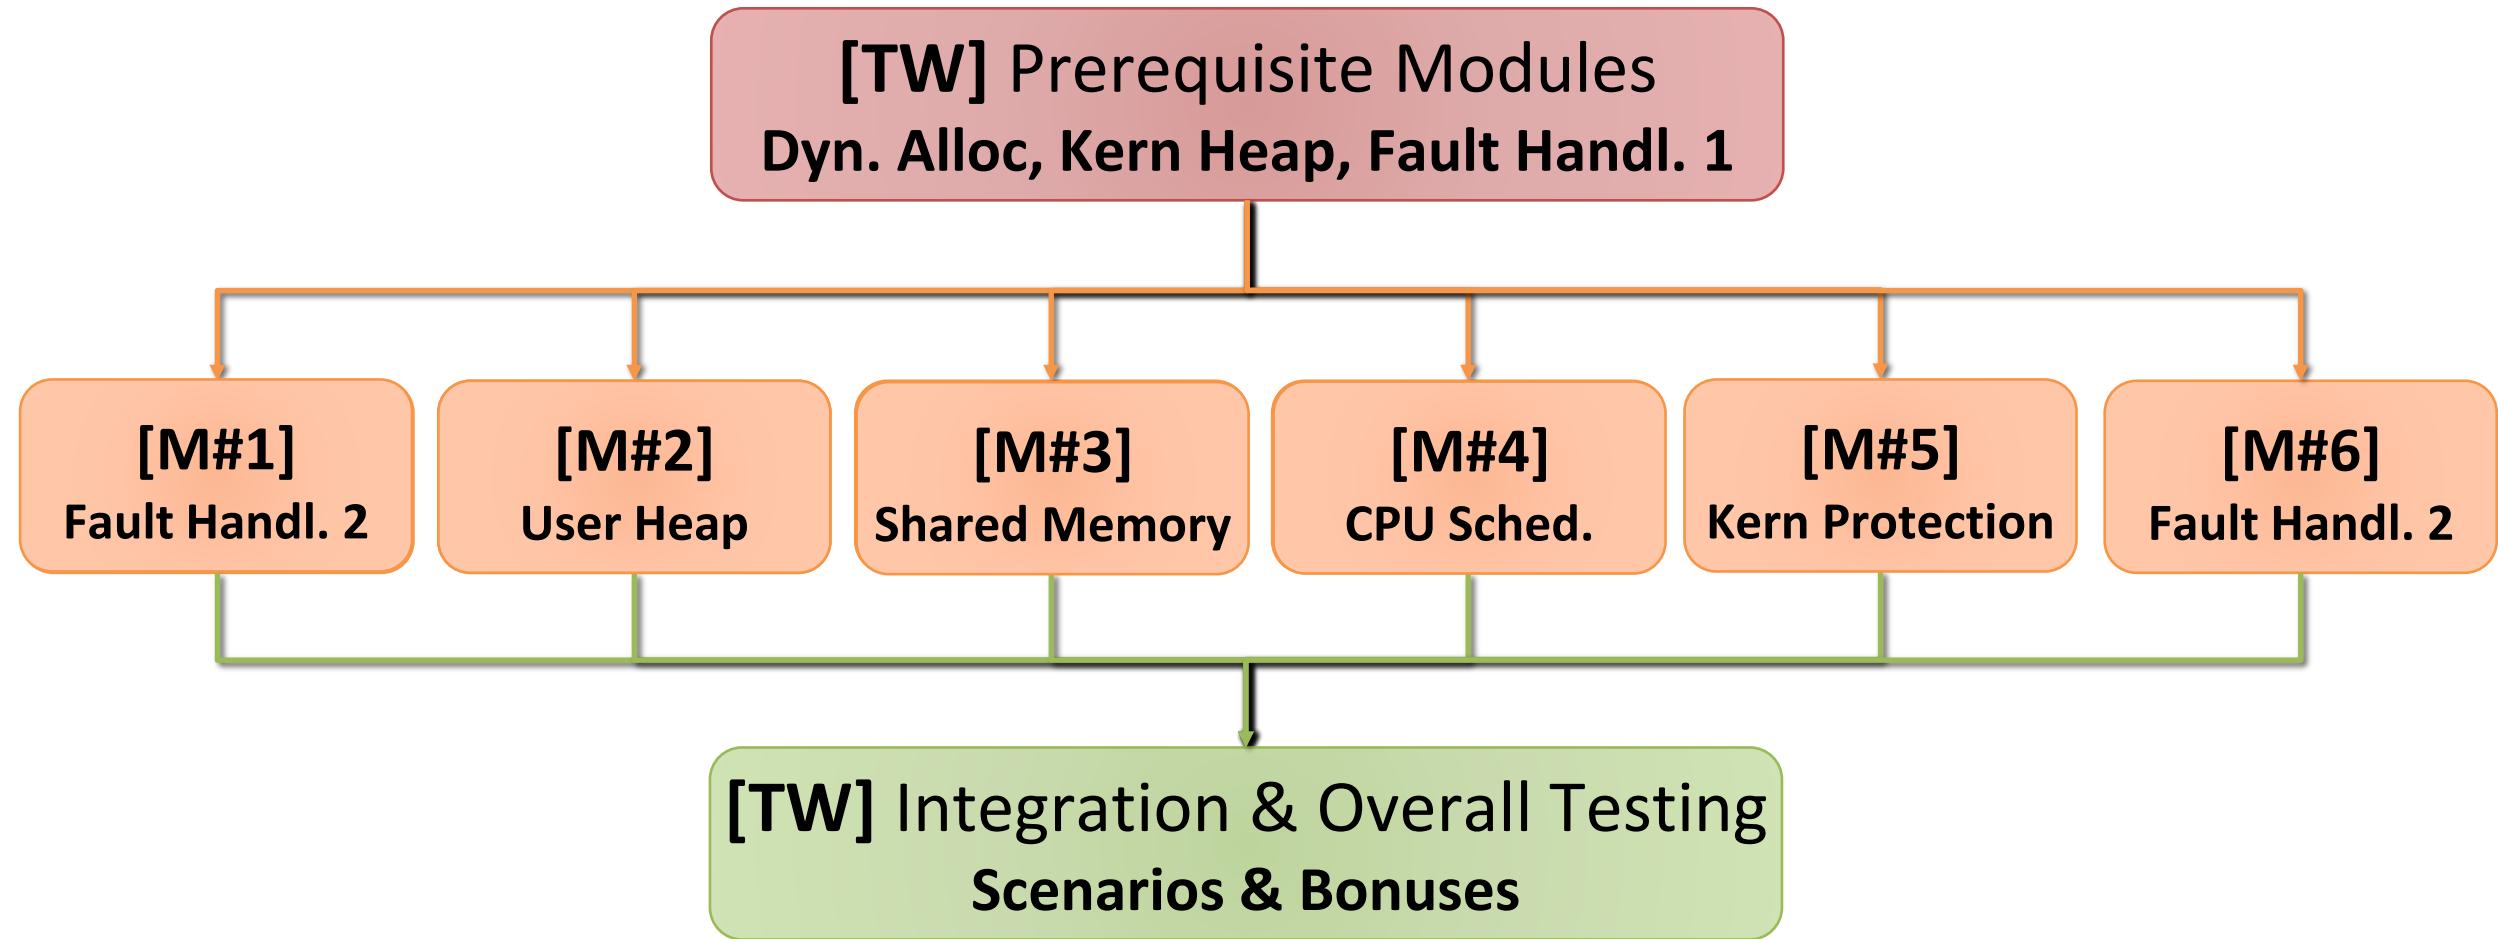
Agenda

- PART 0: ROADMAP
- PART I: PREREQUISITES
- PART II: GROUP MODULES
- PART III: INDIVIDUAL MODULES
- PART IV: OVERALL TESTING & BONUSES

Agenda

- **PART II: GROUP MODULES**
 1. Dynamic Allocator
 2. Kernel Heap
 3. Fault Handler I (Placement)

Project Overview

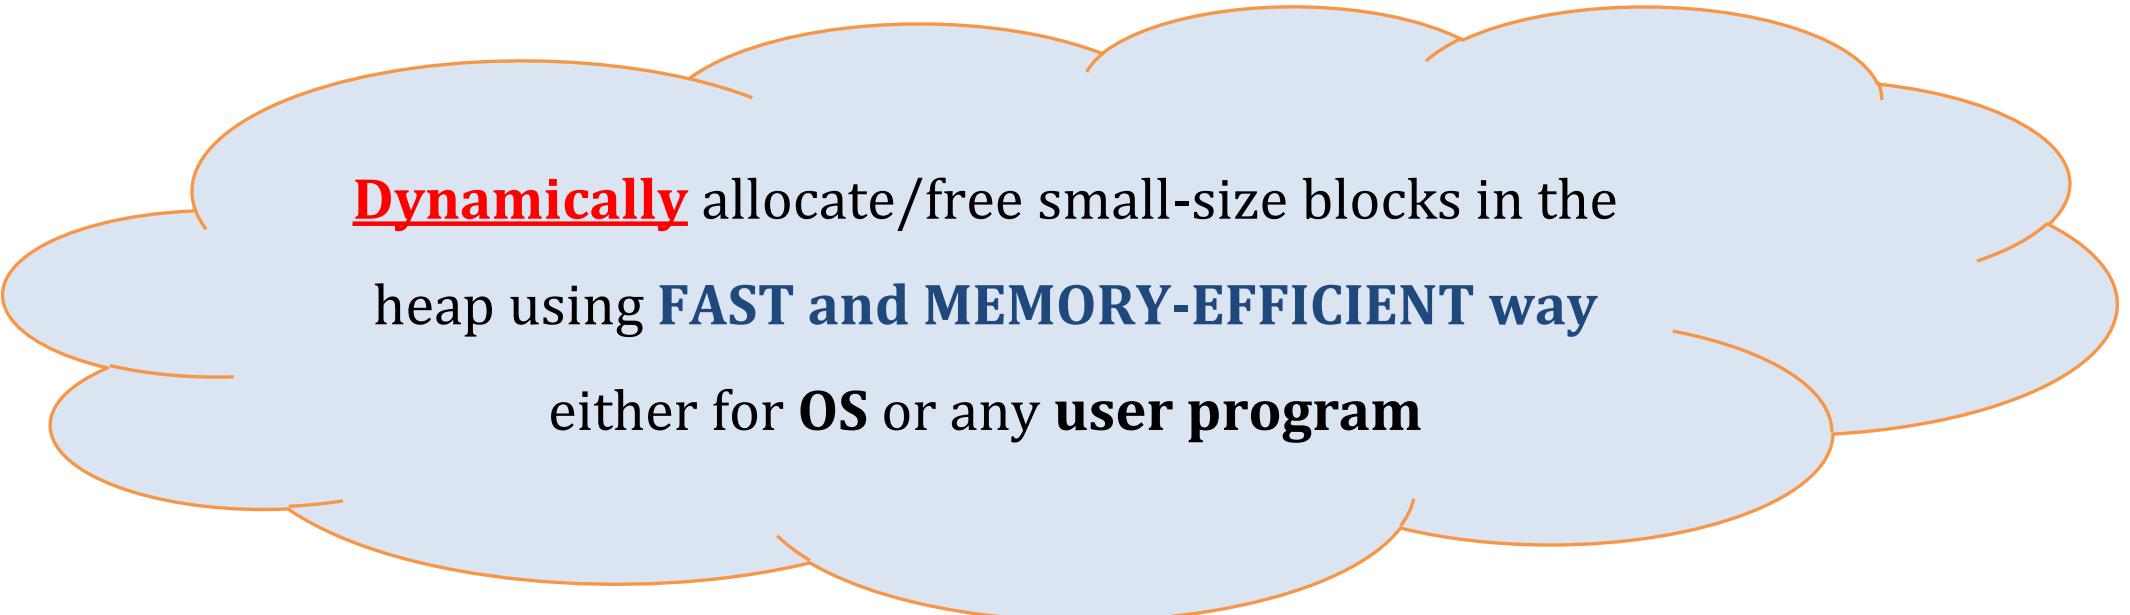




Dynamic Allocator

PART II: GROUP MODULES

Objective



Dynamically allocate/free small-size blocks in the heap using **FAST** and **MEMORY-EFFICIENT** way
either for **OS** or any **user program**

Dynamic Allocator – Functions

The main functions required to handle “Dynamic Allocator” are:

#	Function	File
1	<code>initialize_dynamic_allocator</code>	All essential declarations in: inc/dynamic_allocator.h
2	<code>get_block_size</code>	
3	<code>alloc_block</code>	
4	<code>free_block</code>	Functions definitions <u>TO DO:</u> lib/dynamic_allocator.c
BONUS#1	<code>Block If No Free Block</code>	

Dynamic Allocator – Idea

GOAL: allow the dynamic de/allocation of **small-size blocks** in run-time

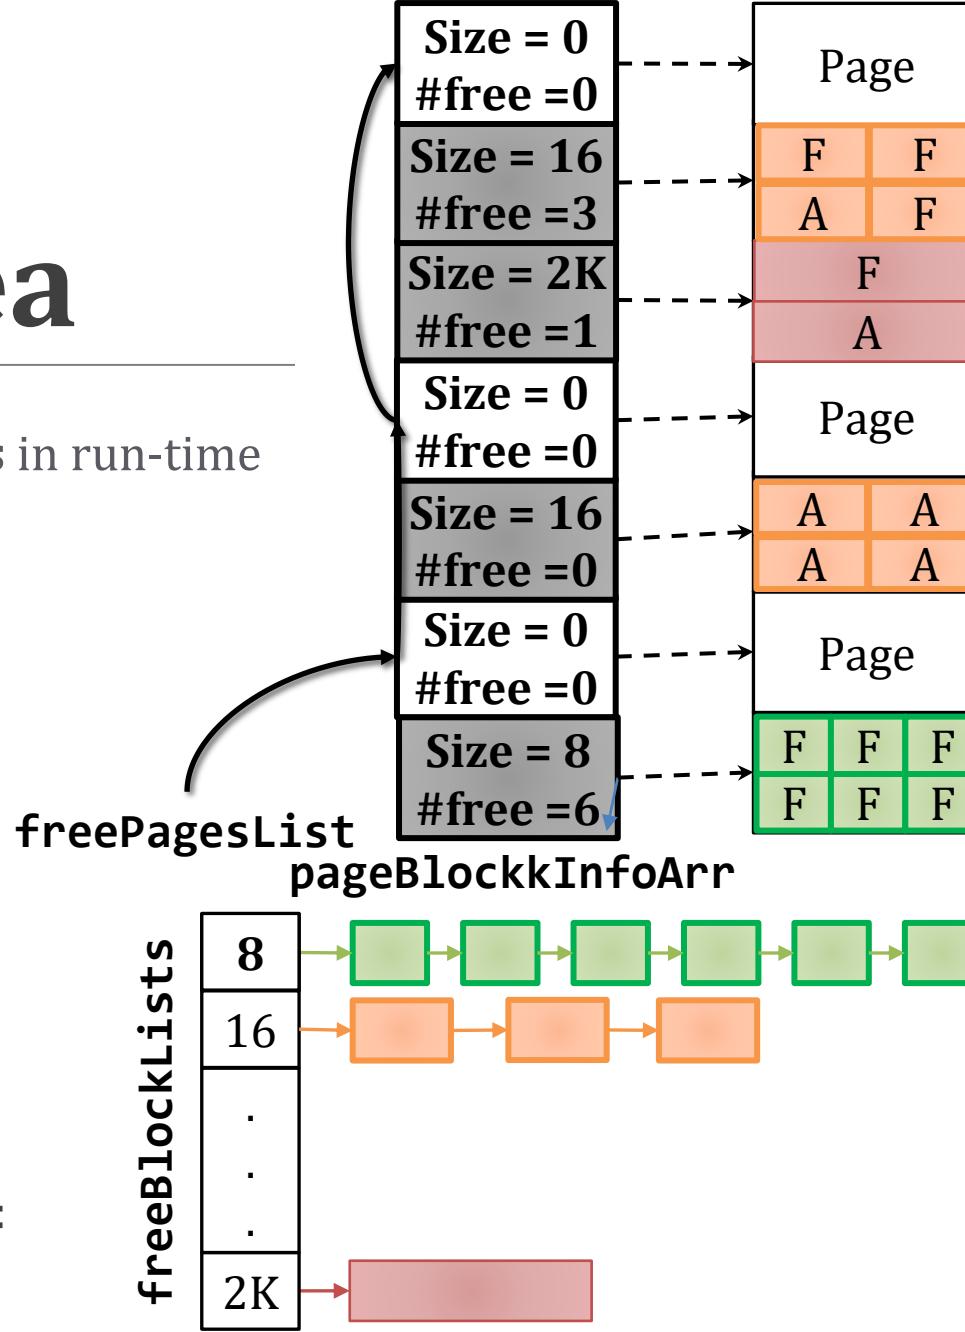
MAIN IDEA:

1. Divide each page into a set of equal-sized blocks
2. Keep track of the free blocks of each size in a linked list
3. For each page in the dynamic allocator space store:
 - a. Size of its blocks
 - b. num of remaining free blocks
4. Keep track of the free pages in another list

Each list is **double-linked** list if elements, each element has:

prev, next Pointers to the adjacent element

Min block size is **8 B...** **WHY?**



Alloc(1.5 KB)

Alloc(20 B)

Alloc(2 KB)

2KB Alloc(55 B)

32B Alloc(62 B)

2KB

2KB

PANIC()

Dynamic Allocator – Allocate

- For the given size, find its nearest **pow-of-2**

1. **CASE1: if a free block exists**

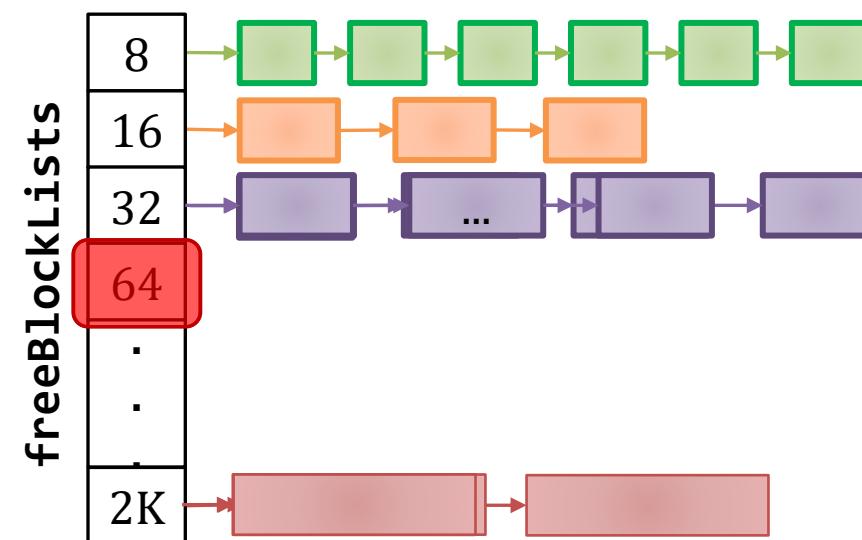
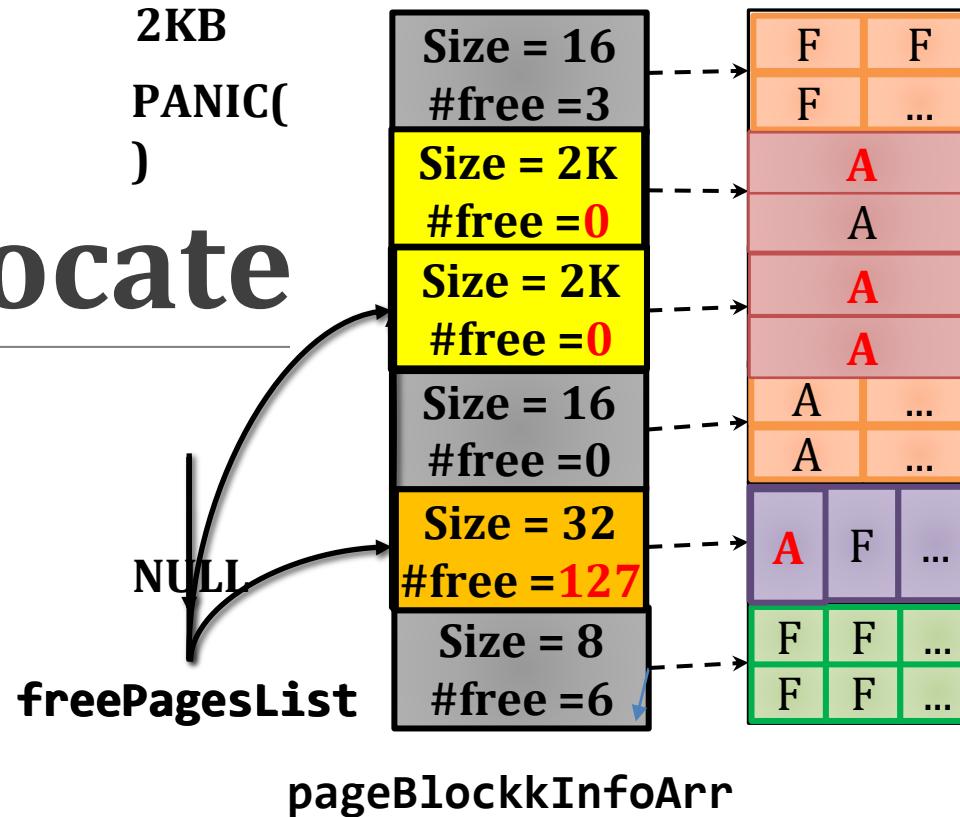
- Allocate it & update the data structures

2. **CASE2: else, if a free page exists**

1. Allocate it from the OS page allocator
2. Split it into blocks of the desired size
3. Add these blocks to the corresponding list
4. Allocate a block & update the data structures

3. **CASE3: else, allocate block from the next list(s)**

4. **CASE4: else, panic(...)**



*X=Alloc(1.5 KB)

*W=Alloc(55 B)

Free(Y)

32B

*Z=Alloc(2 KB)

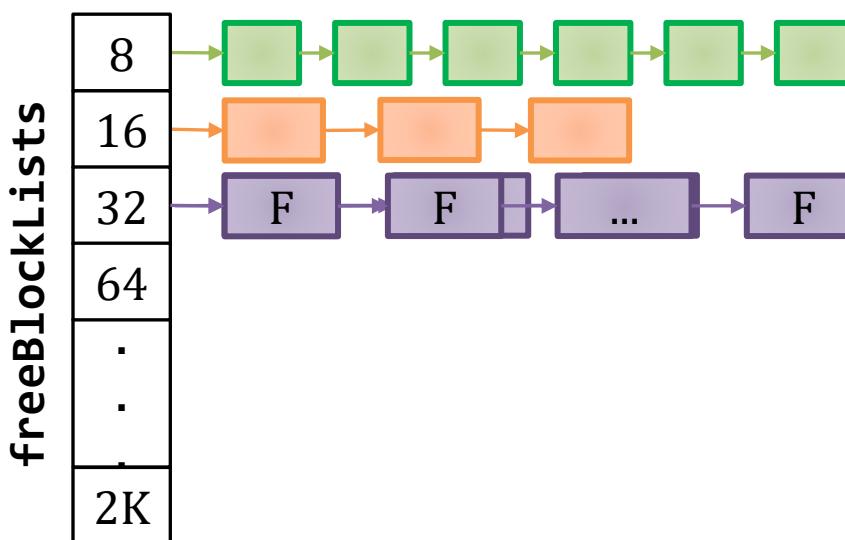
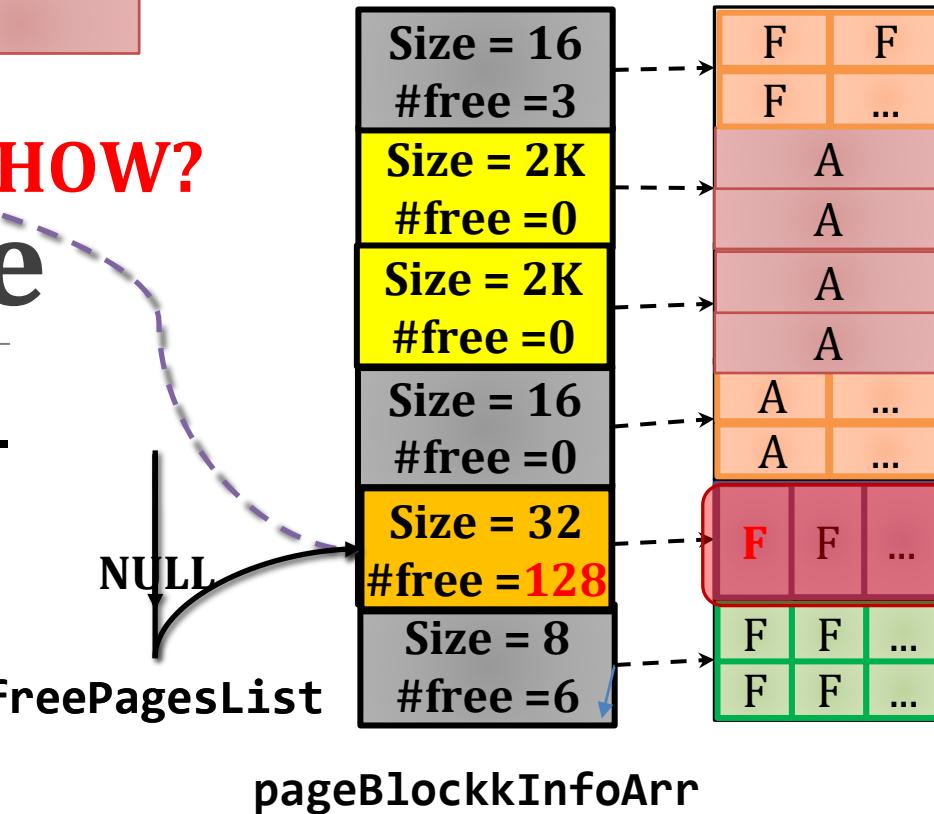
HOW?

Dynamic Allocator - Free

➤ Free the previously allocated block at the given **Virt.**

Addr.

1. Find corresponding size from **pageBlockkInfoArr**
2. Return block to the corresponding list
3. Increment #free blocks in **pageBlockkInfoArr**
4. If entire page becomes free
 1. Remove all **its blocks** from corresponding list in **freeBlockLists**
 2. return it to the free frame list
 3. add it to the **freePagesList**



Dynamic Allocator – Pros & Cons

PROS:

1. **Allocate** has a **BEST case of $O(1)$** and **BOUNDED** by **$O(C)$** ; C is constant = max # blocks/page
2. **Free** is **$O(1)$** and **LINEAR WORST case of $O(K)$** ; K = # free blocks in the corresponding list
3. Can handle **large & small requests efficiently**
4. Allocated page can be returned to the page allocator if becomes free ↗ can **use it for another size**

CONS:

1. **Internal fragmentation** if required size is NOT power-of-two
2. **Bad memory utilization** in case of allocating the request in next level(s)

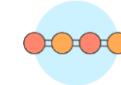
Dynamic Allocator – Given Data

1. Dynamic Allocator Constants

```
#define LOG2_MIN_SIZE (3)                                //3 Bits
#define LOG2_MAX_SIZE (11)                               //11 Bits
#define DYN_ALLOC_MAX_SIZE (32<<20)                  //32 MB
#define DYN_ALLOC_MIN_BLOCK_SIZE (1<<LOG2_MIN_SIZE)    //8 BYTE
#define DYN_ALLOC_MAX_BLOCK_SIZE (1<<LOG2_MAX_SIZE)     //2 KB
```

1. Start & End addresses of the Dynamic Allocator inside the corresponding heap

```
//[3] Limits (to be set in initialize_dynamic_allocator())
uint32 dynAllocStart;
uint32 dynAllocEnd;
```



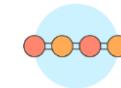
Dynamic Allocator – Given Data

3. Struct containing info about each page with the **prev-next pointer** (to be used as list element)

```
struct PageInfoElement
{
    LIST_ENTRY(PageInfoElement) prev_next_info; /* Linked List Links */
    uint16 block_size;
    uint16 num_of_free_blocks;
};
```

3. Array of Pages Info & Free Pages List

```
LIST_HEAD(PageInfoElement_List, PageInfoElement);
struct PageInfoElement_List freePagesList ;
struct PageInfoElement pageBlockInfoArr[DYN_ALLOC_MAX_SIZE/PAGE_SIZE];
```



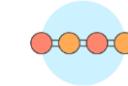
Dynamic Allocator – Given Data

5. Block struct containing the **prev-next pointer** (to be used as list element)

```
struct BlockElement
{
    LIST_ENTRY(BlockElement) prev_next_info; /* linked list links */
};
```

5. Free Block List of ALL sizes that holds elements of type **BlockElement**

```
LIST_HEAD(BlockElement_List, BlockElement);
struct BlockElement_List freeBlockLists[LOG2_MAX_SIZE - LOG2_MIN_SIZE + 1] ;
```



Dynamic Allocator - Given Functions

1. Get start VA of the page from the corresponding Page Info pointer

```
uint32 to_page_va(struct PageInfoElement* ptrPageInfo);
```

2. Get a page from the Kernel Page Allocator for DA (i.e. Allocate it) (either for **USER** or **KERNEL**)

```
int get_page(void* va);
```

In **KERNEL**: it calls `alloc_page()` in `paging_helpers.c` (feel free to edit it if necessarily)

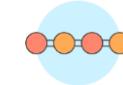
In **USER**: it calls `_sys_allocate_page()` which, in turn, calls also `alloc_page()`

3. Return a page to the Kernel Page Allocator (i.e. Free It) (either for **USER** or **KERNEL**)

```
void return_page(void* va);
```

In **KERNEL**: it calls `unmap_frame()`.

In **USER**: it calls `_sys_unmap_frame()` which, in turn, calls also `unmap_frame()`

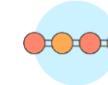


#1: Initialize

```
void initialize_dynamic_allocator(uint32 daStart, uint32  
                                 daEnd);
```

Description:

- Initialize the dynamic allocator starting from the given “daStart” to “daEnd”
- The following items should be initialized here:
 1. DA Limits
 2. Array of Page Info
 3. Free Page List
 4. Free Block Lists

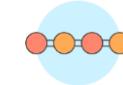


#2: Get Block Size

```
uint32 get_block_size(void *va)
```

Description:

- Get the block size at the given VA from its corresponding entry in the **pageBlockInfoArr**

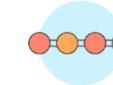


#3: Allocate

```
void *alloc_block(uint32 required_size)
```

Description:

- Allocate new block with the given size
- For the given size, **find** its **nearest** pow-of-two
- **Allocate** new block according to the previously explained cases
- **Return** the start address of the allocated block OR NULL if the requested size is 0

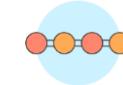


#4: Free

```
void free_block(void* va)
```

Description:

- **Free** the previously allocated block at the given address “**va**”
- **Follow the previously explained steps**



BONUS#1: Block If No Free Block

```
void *alloc_block(uint32 required_size)
```

Description:

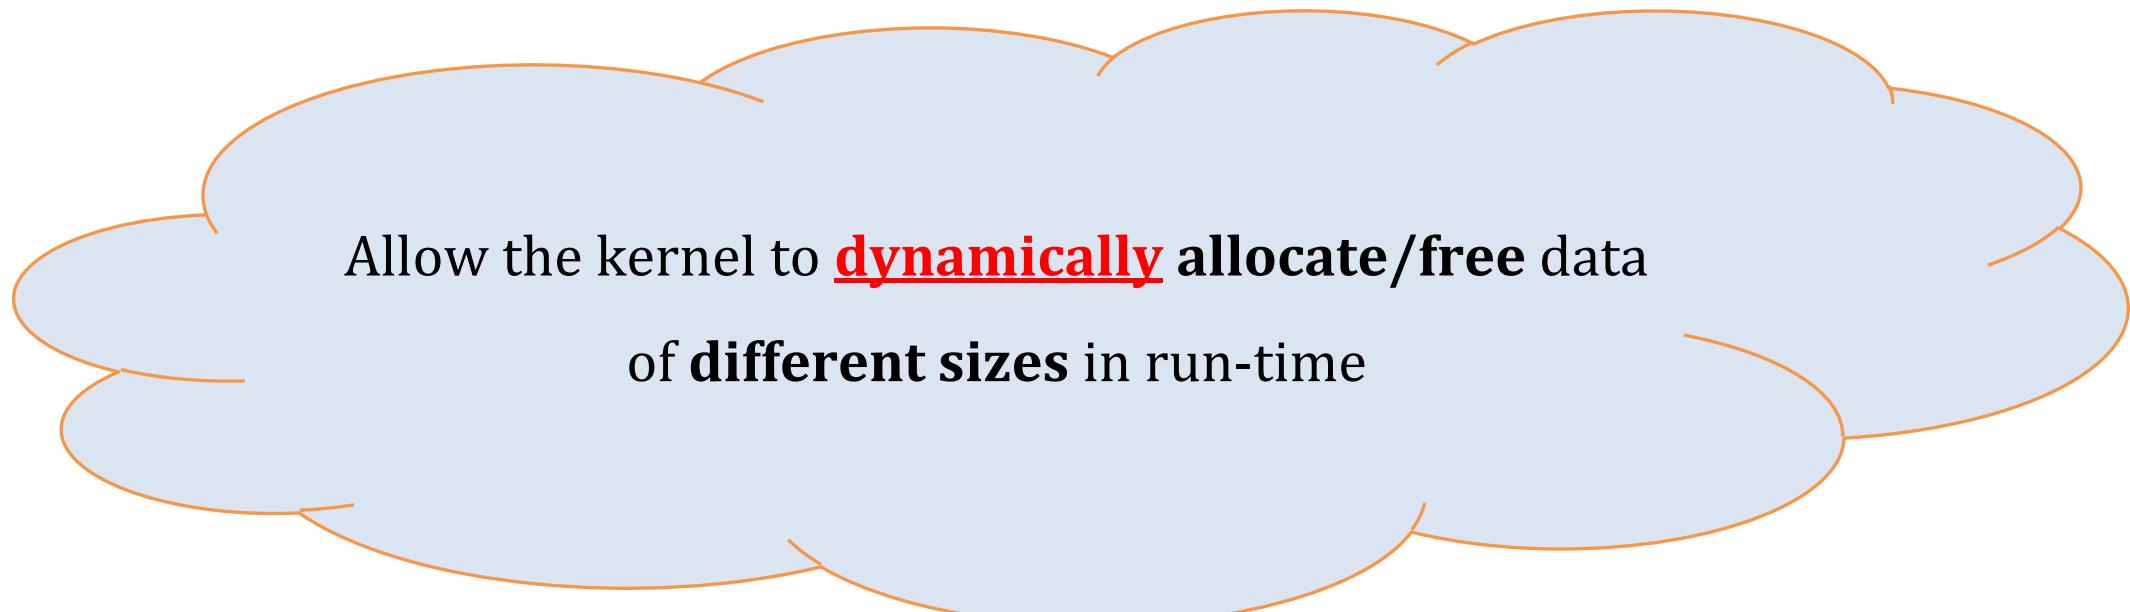
- If there's no free block, instead of panic, **BLOCK** the process until a block becomes available



Kernel Heap

PART II: GROUP MODULES

Objective



Kernel Heap

IMPORTANT NOTES

Enable the KERNEL HEAP: '[inc/memlayout.h](#)' set **USE_KHEAP** by 1
This module is **ESSENTIAL** for all other modules

CAUTION

During your solution, any SHARED data MUST be PROTECTED
by critical section via LOCKS

Kernel Heap – Functions

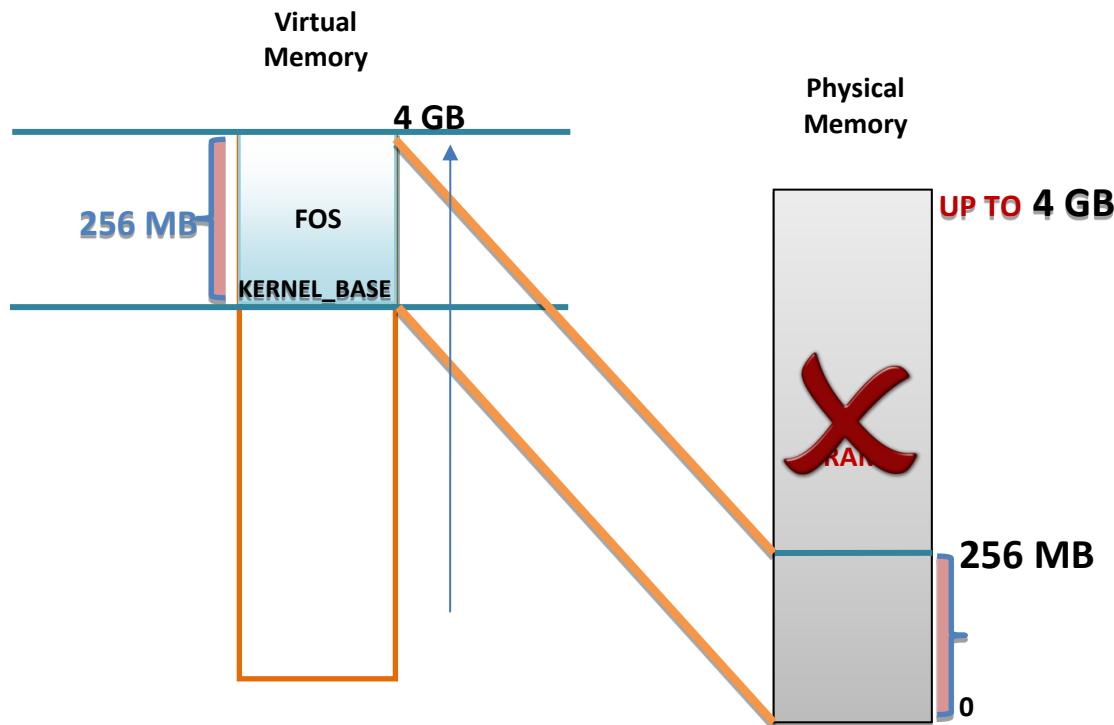
The main functions required to handle “Kernel Heap” are:

#	Function	File
1	<code>kmalloc</code>	All essential declarations in: Kern/mem/kheap.h
2	<code>kfree</code>	
3	<code>kheap_virtual_address</code>	
4	<code>kheap_physical_address</code>	Functions definitions <u>TO DO</u> in: Kern/mem/kheap.c
BONUS 1	<code>krealloc</code>	
BONUS 2	<code>Fast Page Allocator</code>	

Kernel Heap – What is new?

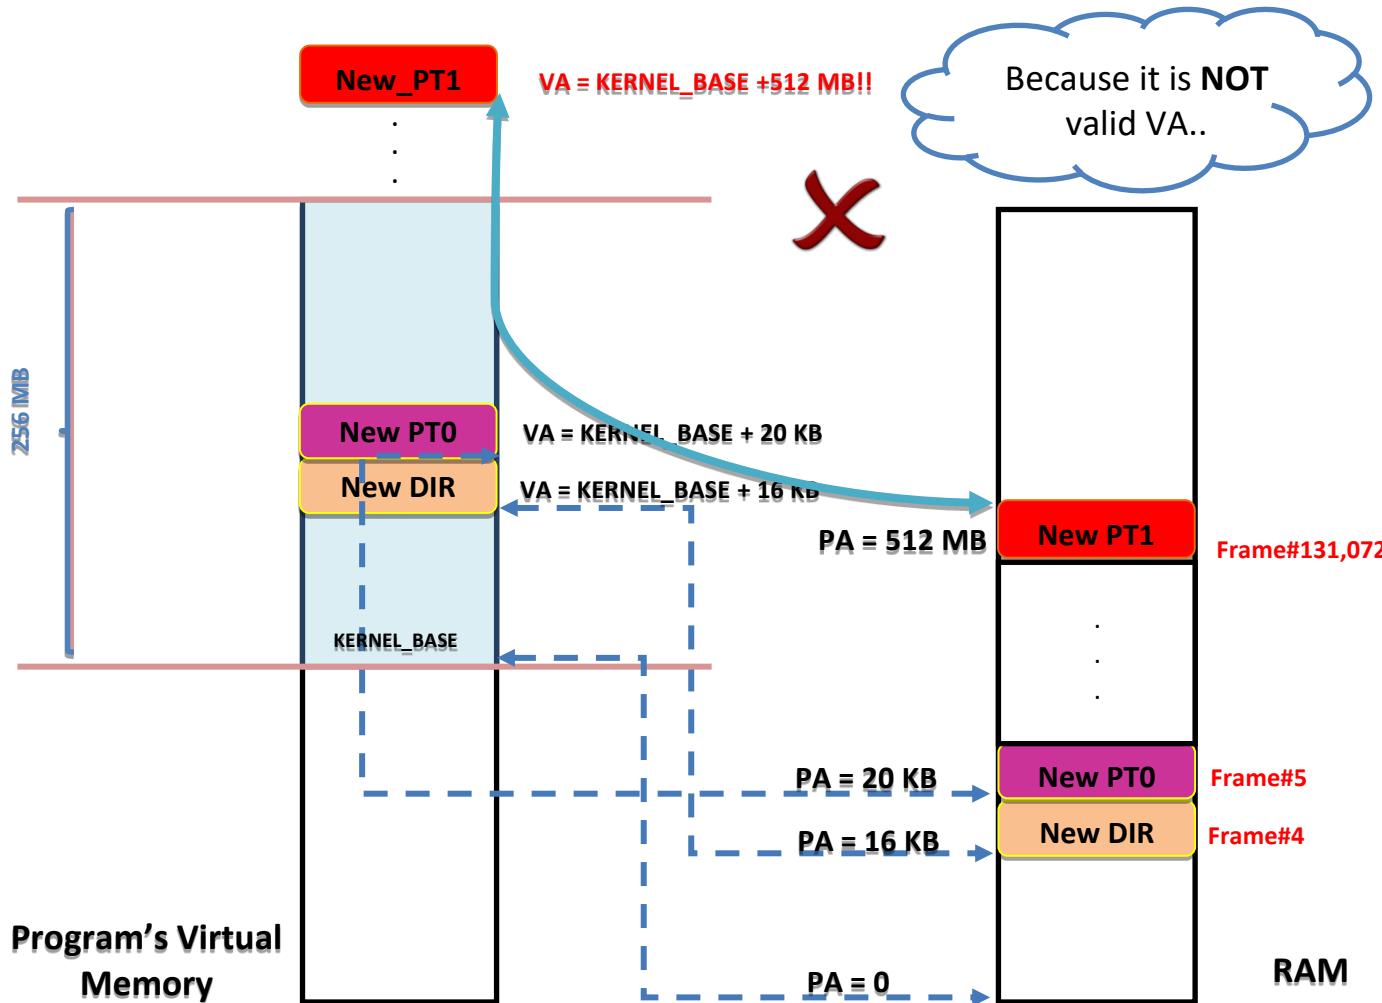
Current: Kernel is **one-to-one** mapped to 256 MB RAM

Problem: Kernel can't directly access beyond 256 MB RAM



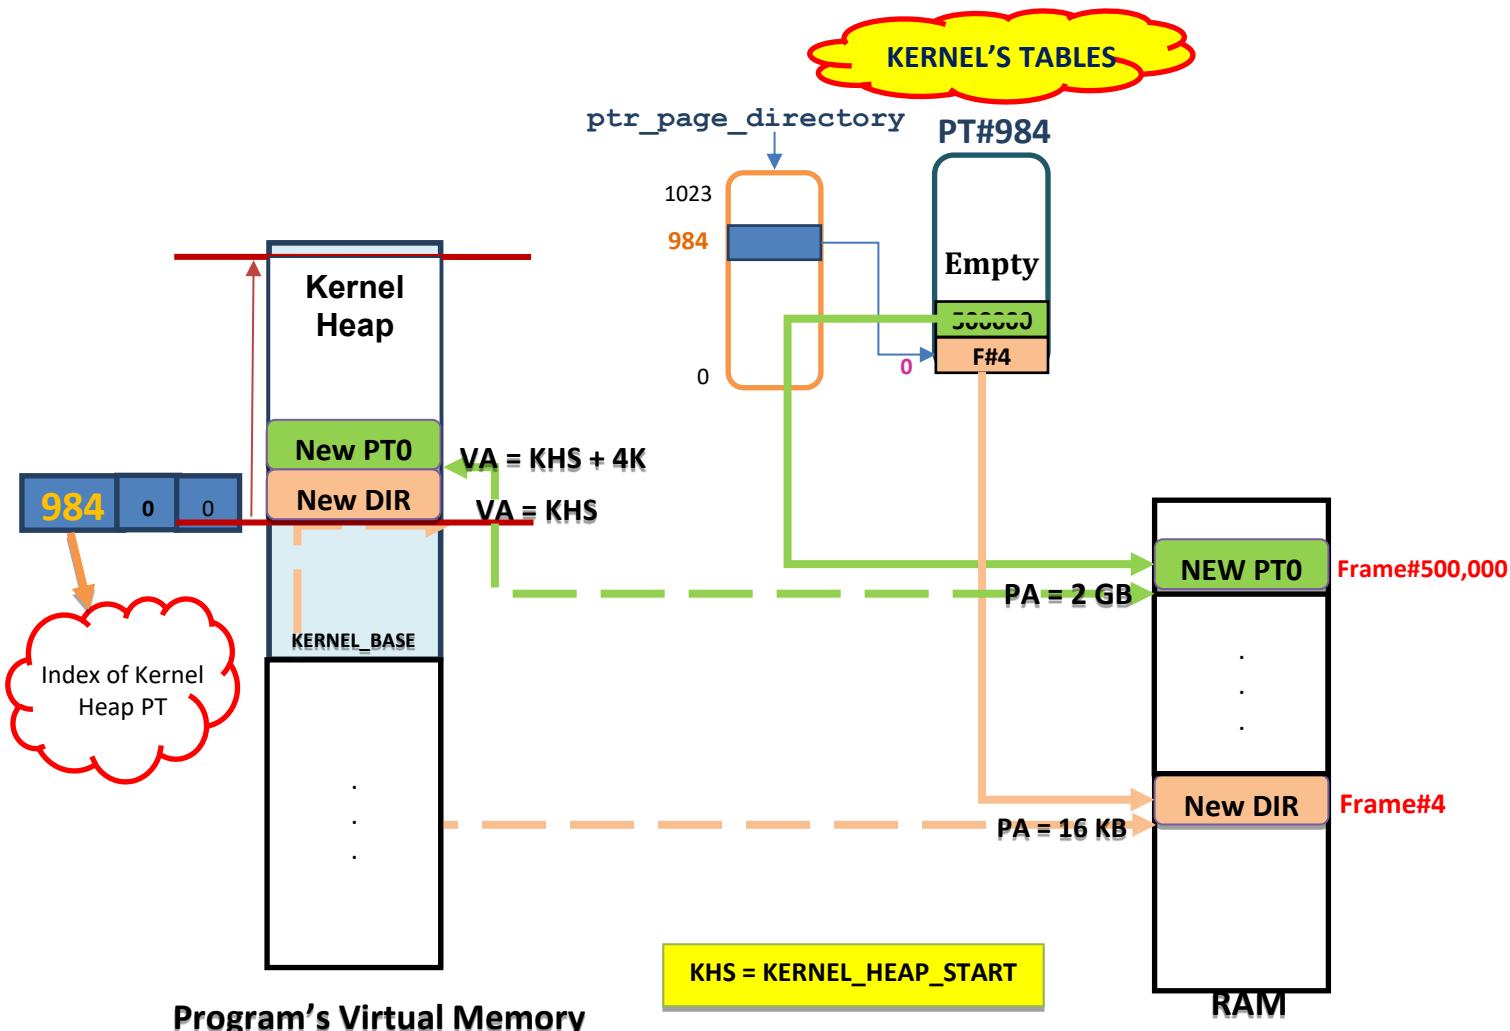
Kernel Heap – What is new?

- Example: Kernel can't directly access beyond 256 MB RAM



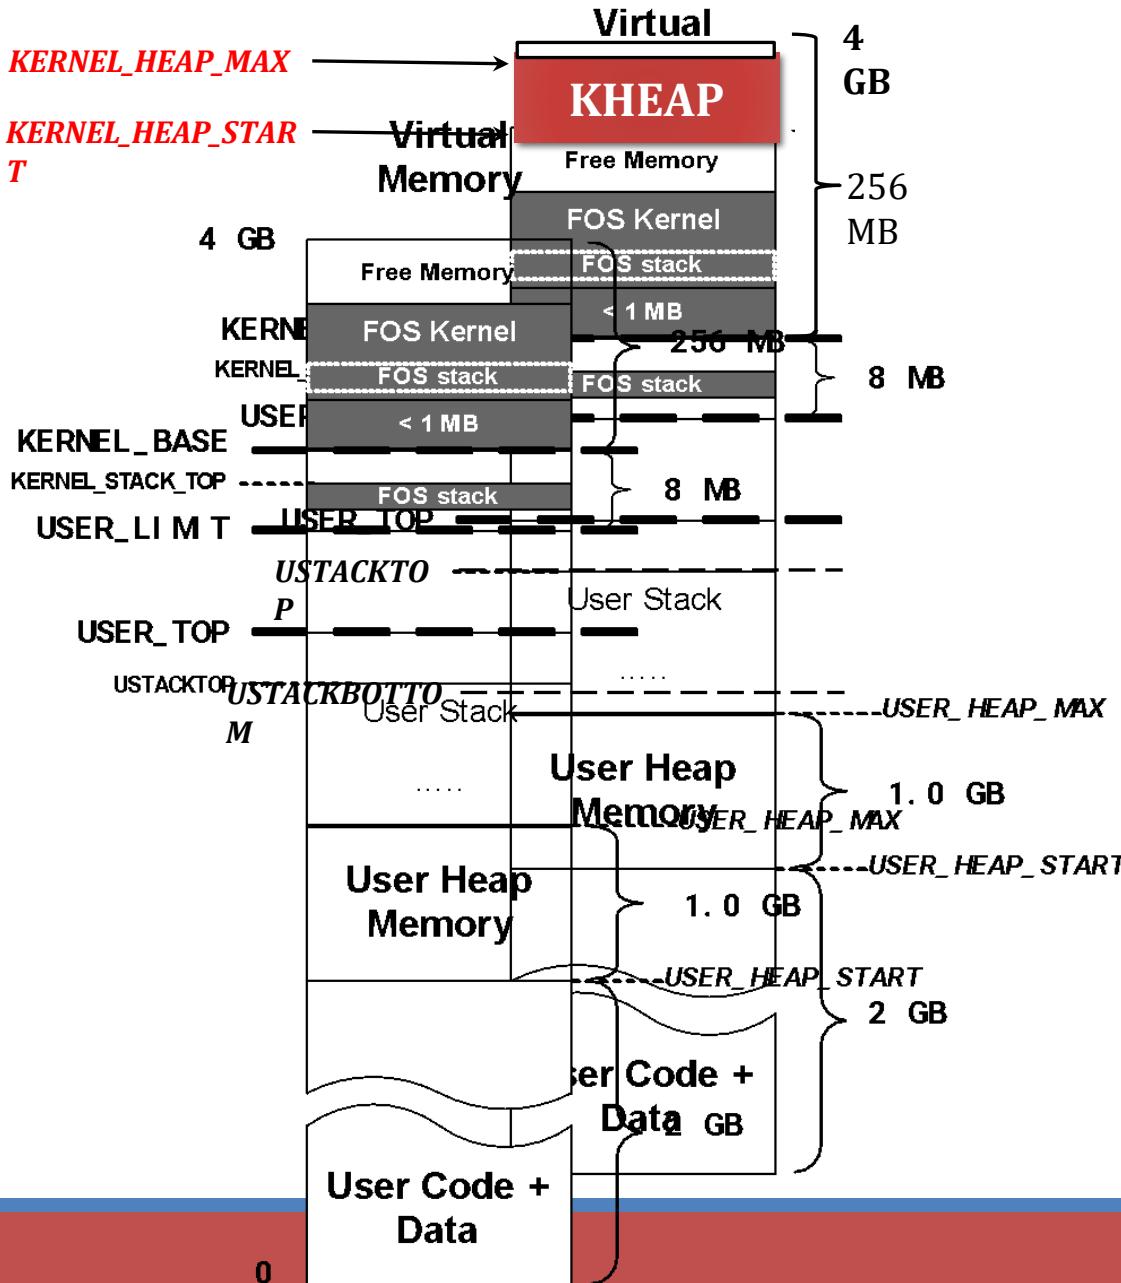
Kernel Heap – What is new?

- Solution: Kernel Heap for dynamic allocations (**No 1-1 map**)



Kernel Heap – What is new?

Kernel Heap lies at the
end of the virtual space



Kernel Heap – Allocation Types?

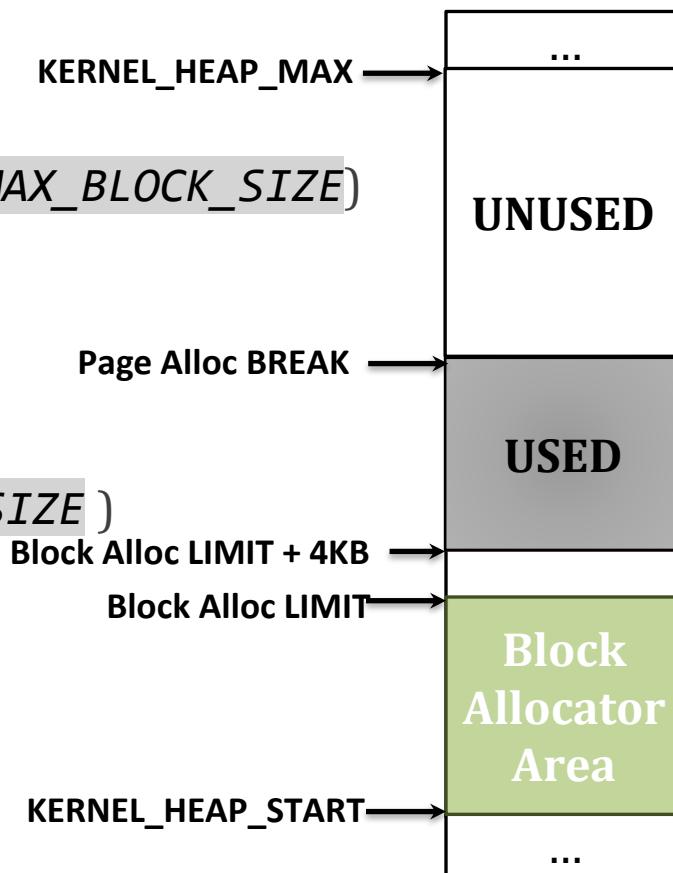
There're **TWO** types of allocator

1. Block Allocator

1. Used to allocate **small blocks** (with size **LESS OR EQUAL** `DYN_ALLOC_MAX_BLOCK_SIZE`)
2. Use Dynamic Allocator Functions
3. Range: `[KERNEL_HEAP_START, BLK ALLOC LIMIT]`

2. Page Allocator

1. Used to allocate **chunk of pages** (with size **>** `DYN_ALLOC_MAX_BLOCK_SIZE`)
2. Allocation is done on **page boundaries** (i.e. internal fragmentation)
3. Range: `[BLK ALLOC LIMIT + PAGE_SIZE, KERNEL_HEAP_MAX]`
 1. **USED** Area: `[BLK ALLOC LIMIT + PAGE_SIZE, PAGE ALLOC BREAK]`
 2. **UNUSED** Area: `[PAGE ALLOC BREAK, KERNEL_HEAP_MAX]`

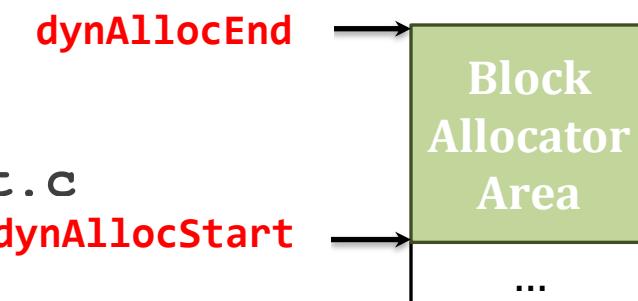


Kernel Heap – Block Allocator

1. Has 2 limits:
 1. `dynAllocStart`: begin of block allocator area
 2. `dynAllocEnd`: end of block allocator area
2. Use Dynamic Allocator with its data structure
3. **Already initialized**, together with the dynamic allocator itself inside:

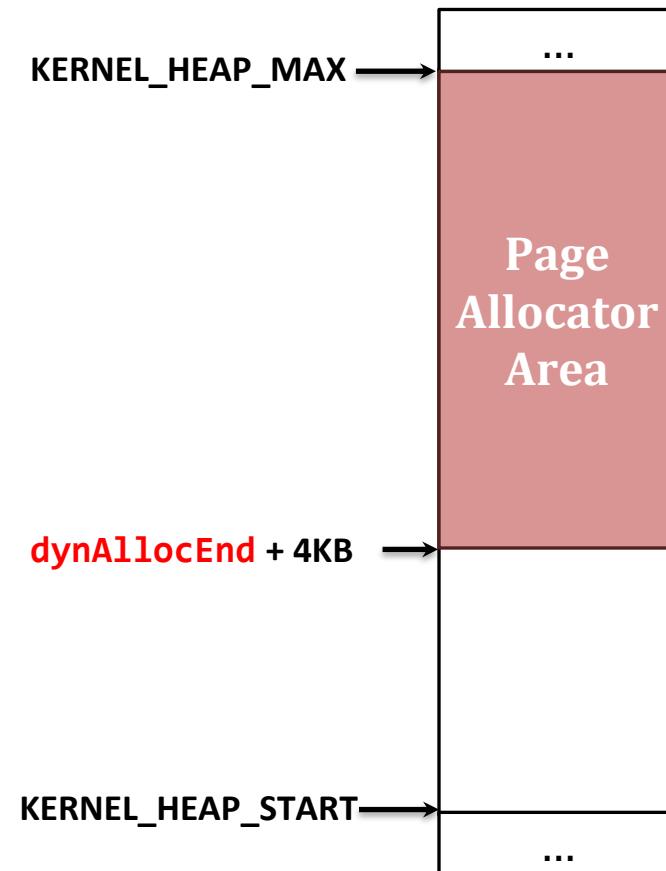
`int kheap_init(...)` defined in `kern/mem/kheap.c`

- This function, in turn, is **already called** in `FOS_initialize()` in `init.c`



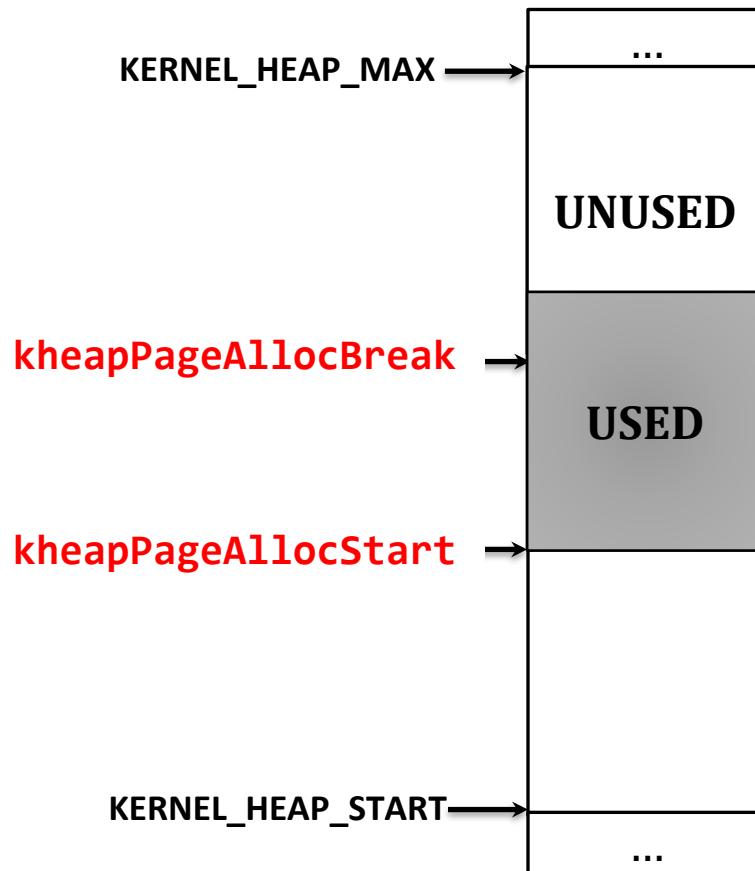
Kernel Heap – Page Allocator

- Should start at **one-page after the block allocator limit**
- Allocation is done on **page boundaries** (multiple of 4KB)
 - i.e. **internal fragmentation** can occur
- All required pages should be **allocated & mapped** by OS
- Allocation Strategy: **CUSTOM FIT**



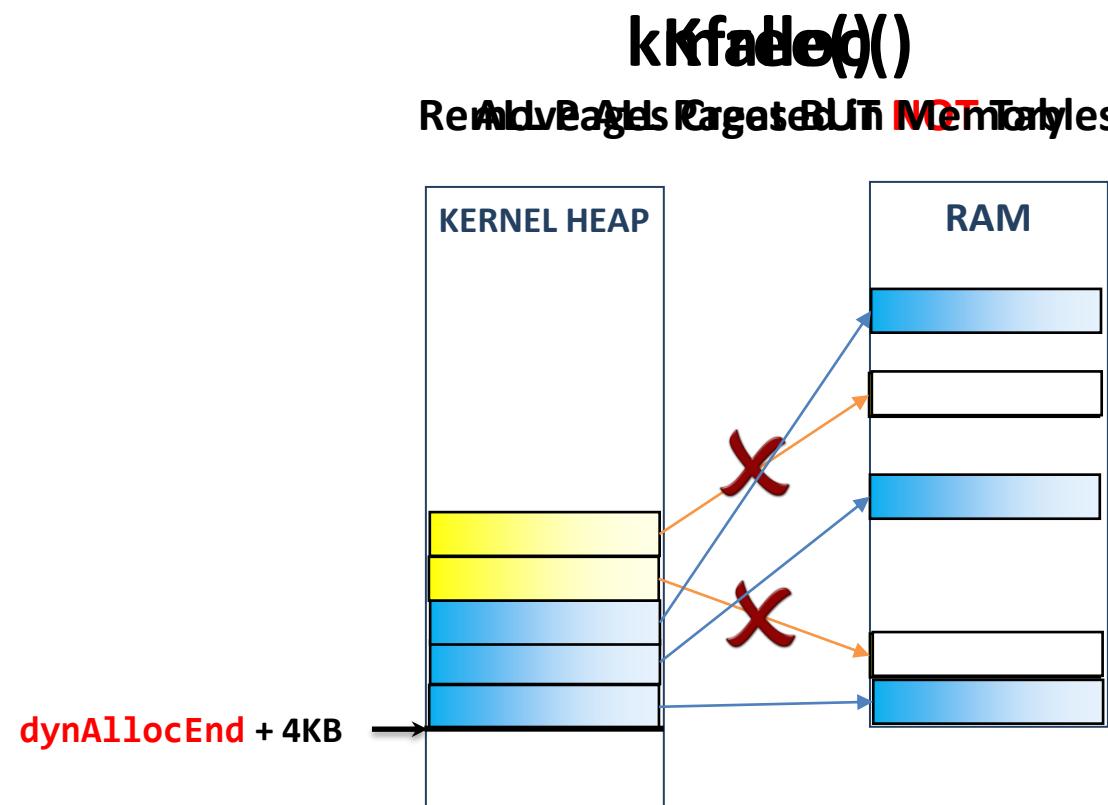
Kernel Heap – Page Allocator

1. Has 2 limits defined in **kern/mem/kheap.h**:
 1. **kheapPageAllocStart**: begin of page allocator area
 2. **kheapPageAllocBreak**: end of currently used area
2. malloc/free can move **kheapPageAllocBreak** up/down



Kernel Heap – Page Allocator

1. `kmalloc()`: dynamically allocate space
2. `kfree()`: delete a previously allocated space



Kernel Heap – Page Allocator (kmalloc)

Allocate pages on 4KB granularity

`ptr3 = kmalloc (3`

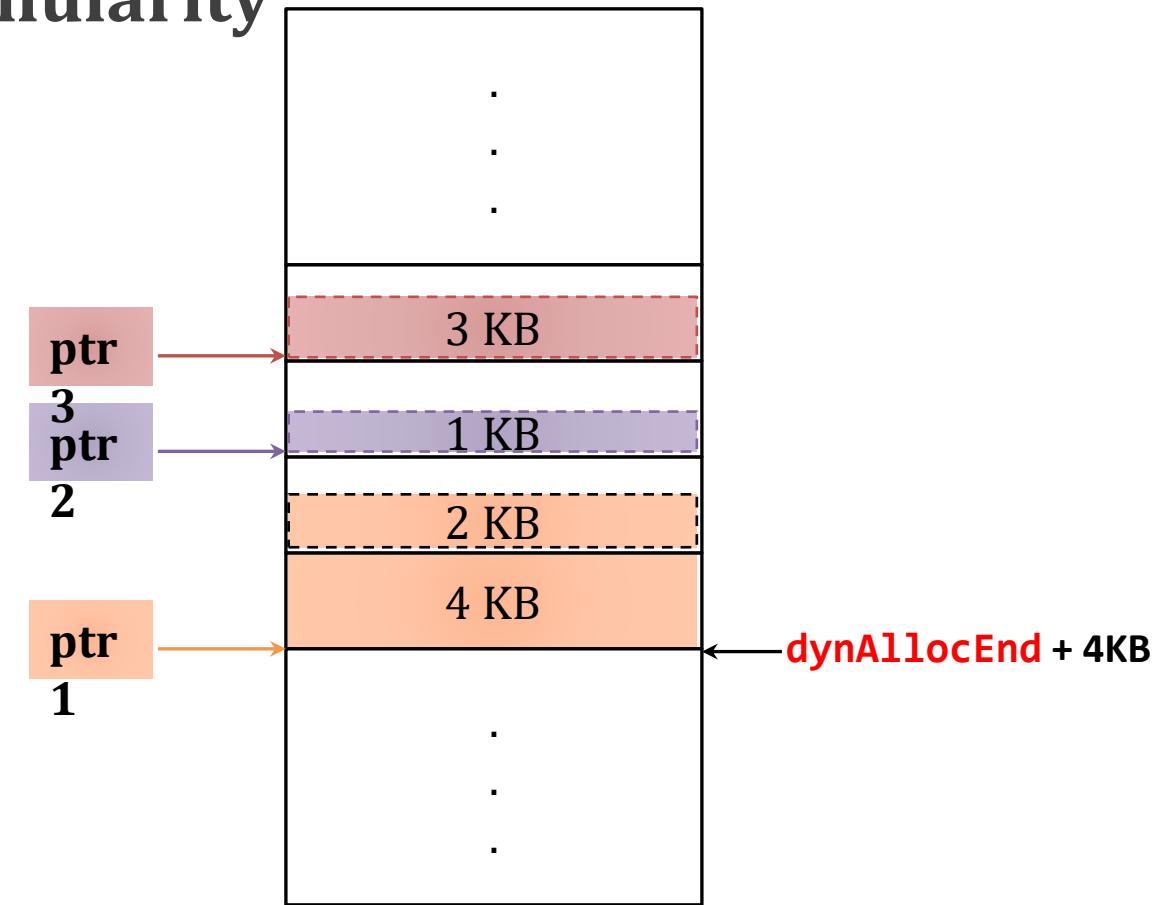
KB)

`ptr2 = kmalloc (1`

KB)

`ptr1 = kmalloc (6`

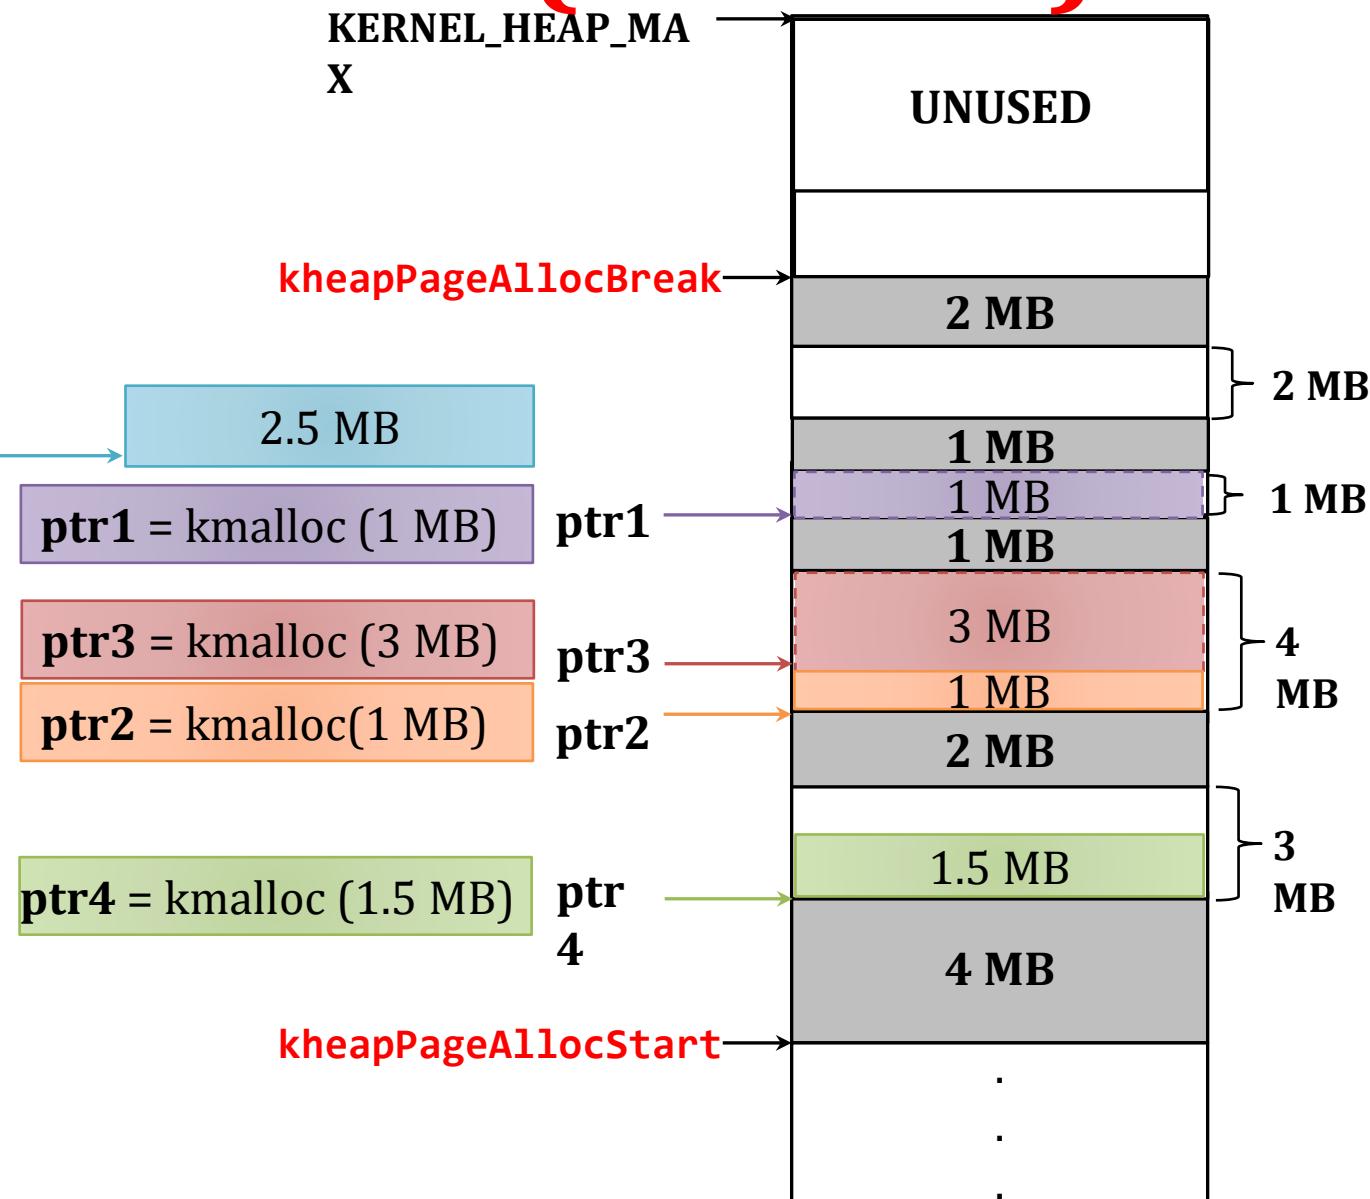
KB)



Kernel Heap – Page Allocator (kmalloc)

CUSTOM FIT Strategy

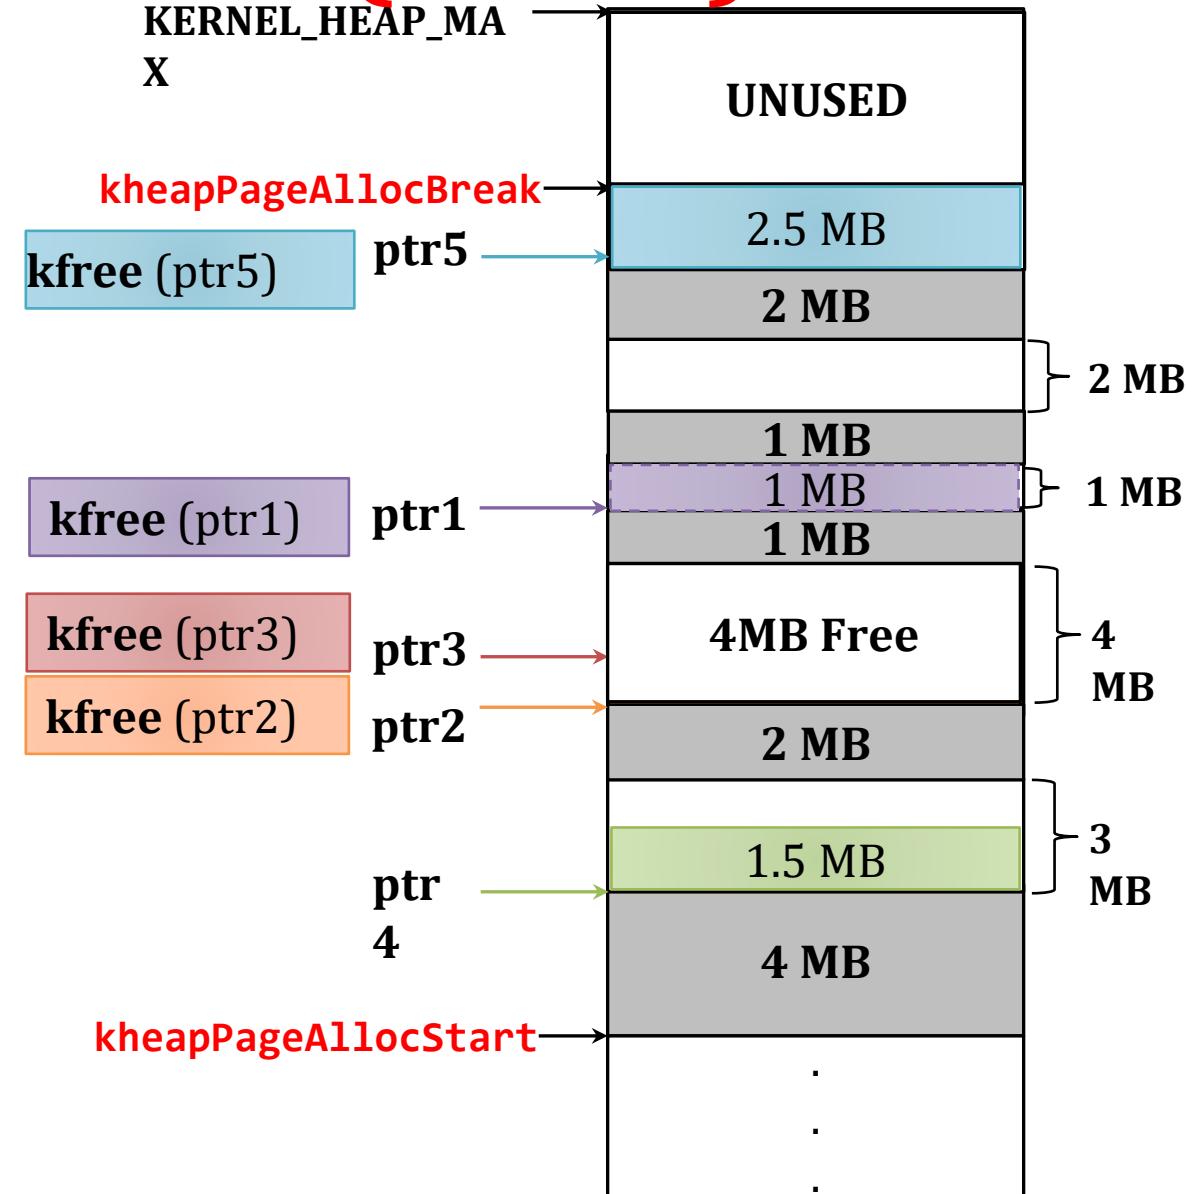
1. Search for **EXACT** fit
2. if not found, **ptr5 = kmalloc (2.5 MB)** search for **WORST** fit till break
3. if not found, extend **BREAK** if available
4. if not available, return **NULL**



Kernel Heap – Page Allocator (kfree)

Make sure to:

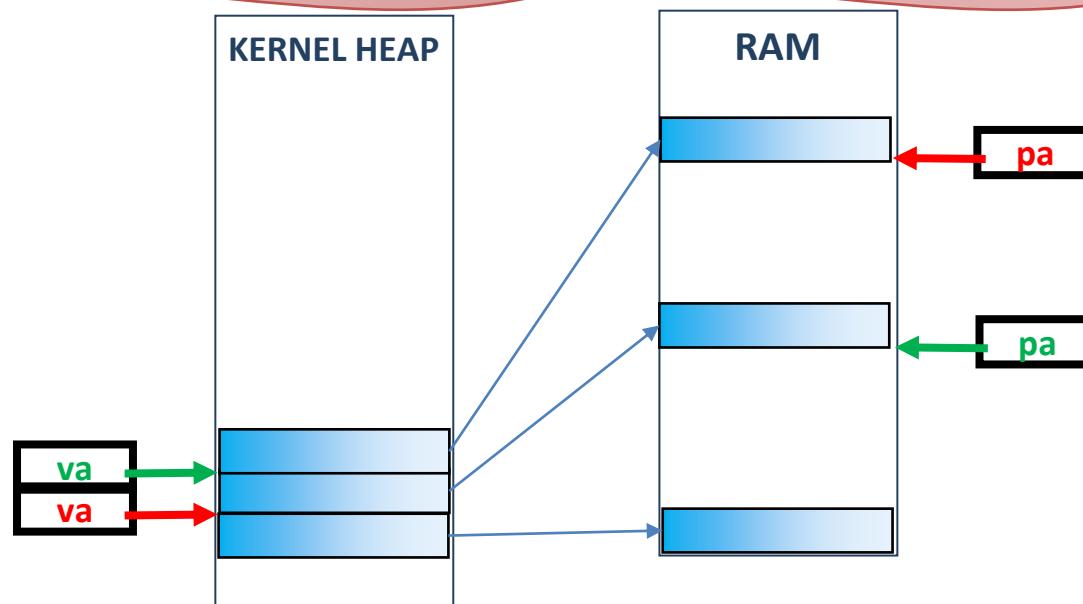
1. Merge adjacent blocks
2. Update the **kheapPageAllocBreak** if freeing the last space



Kernel Heap – Page Allocator

3. `kheap_physical_address()`: find physical address of the given kernel one
4. `kheap_virtual_address()`: find kernel virtual addr of the given physical one

IMP: Both MUST be Efficient ~O(1)



#5: kmalloc()

```
void* kmalloc(unsigned int size)
```

Description:

1. If $\text{size} \leq DYN_ALLOC_MAX_BLOCK_SIZE$: **[BLOCK ALLOCATOR]**
 - Use dynamic allocator to allocate the required space
2. Else: **[PAGE ALLOCATOR]**
 - Allocate & map the required space on page-boundaries using CUSTOM FIT strategy
 - If failed to allocate: return NULL

#6: kfree()

```
void kfree(void* virtual_address)
```

Description:

1. If virtual address inside the **[BLOCK ALLOCATOR]** range
 - Use dynamic allocator to free the given address
2. If virtual address inside the **[PAGE ALLOCATOR]** range
 - FREE the space of the given address from RAM
3. Else (i.e. invalid address): should **panic** (...)

#7: kheap_physical_address()

```
unsigned int kheap_physical_address(unsigned int virtual_address)
```

Description:

1. return the physical address corresponding to given virtual_address (**including offset**)
2. If no mapping, return 0.
3. It should work for both **[BLOCK ALLOCATOR]** and **[PAGE ALLOCATOR]**
4. It should run in **O(1)**

#8: kheap_virtual_address()

```
unsigned int kheap_virtual_address(unsigned int physical_address)
```

Description:

1. return the virtual address corresponding to given physical_address (**including offset**)
2. If no mapping, return 0.
3. It should work for both **[BLOCK ALLOCATOR]** and **[PAGE ALLOCATOR]**
4. It should run in **O(1)**

BONUS#1: krealloc()

```
void *krealloc(void *virtual_address, uint32 new_size)
```

Description:

1. Attempts to resize the allocated space at given virtual address to "**new size**" bytes, possibly moving it in the heap.
2. If **successful**, returns the **new virtual address**.
3. On **failure**, returns a **NULL** pointer, and the old virtual address remains valid.
4. A call with `virtual_address = null` is equivalent to `kmalloc()`
5. A call with `new_size = zero` is equivalent to `kfree()`
6. It should work in **[BLOCK ALLOCATOR]** , **[PAGE ALLOCATOR]** and **CROSS-ALLOCATORS**

BONUS#2: Fast Page Allocator

Description:

EFFICIENT implementation of the **Page Allocator** using **suitable data structures**



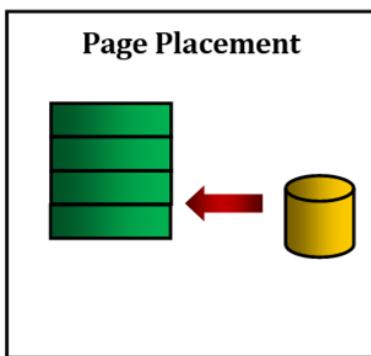
Fault Handler I

(Placement)

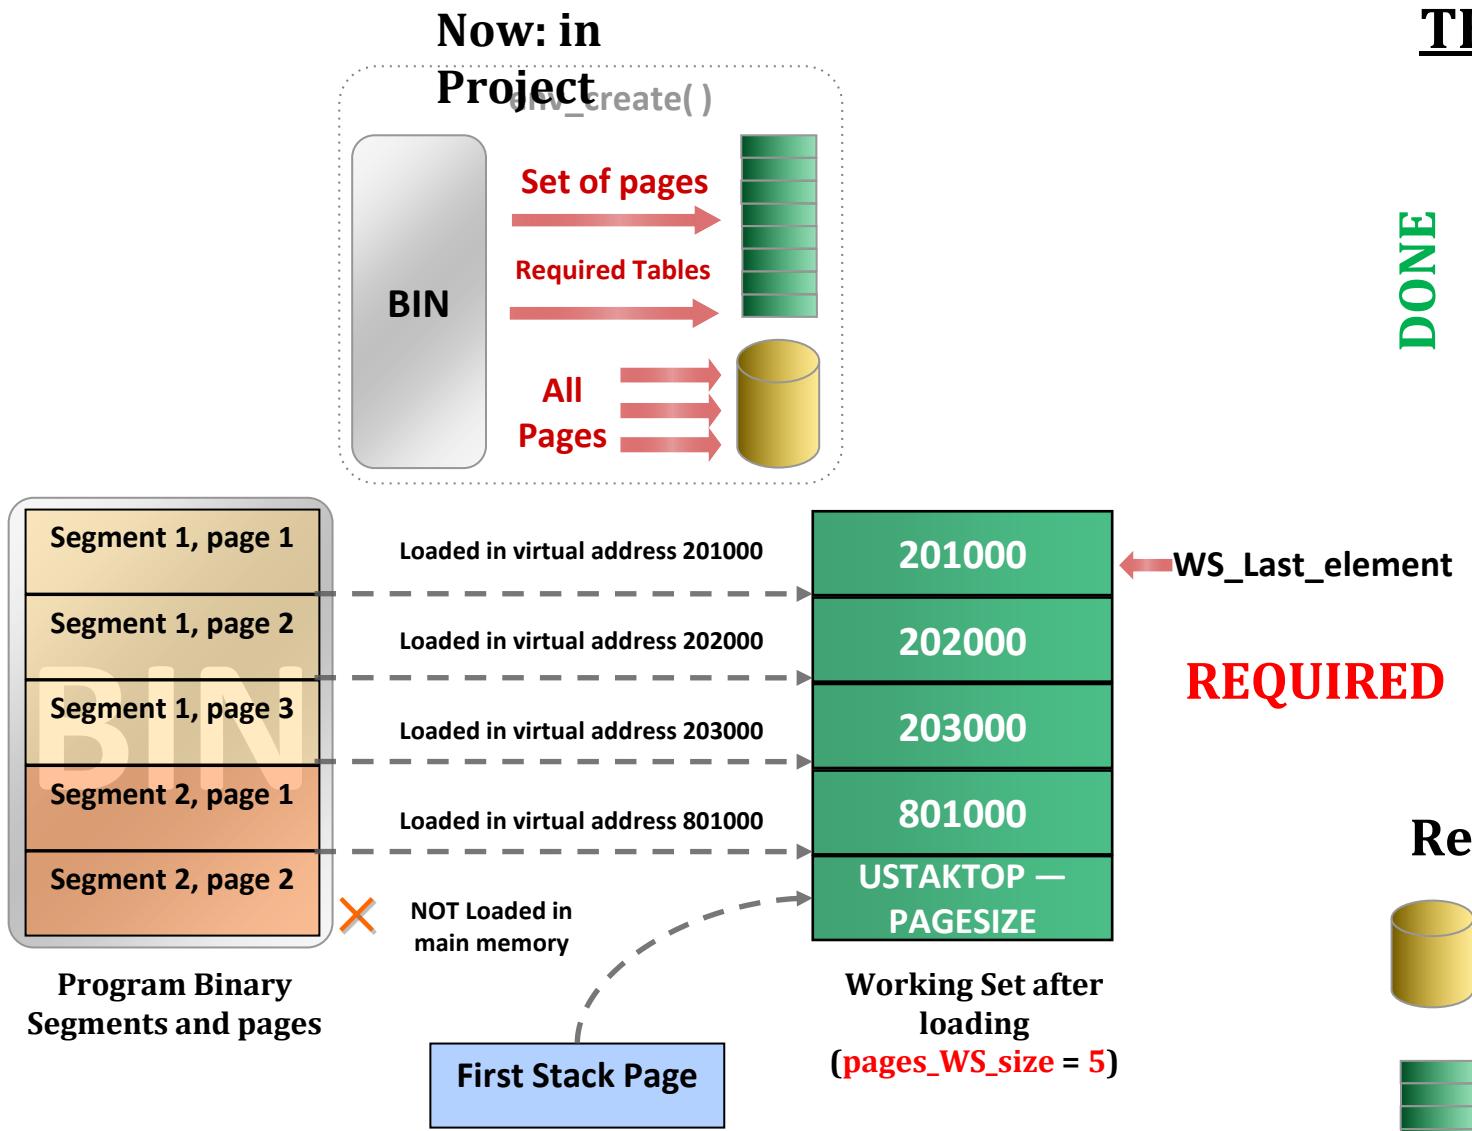
PART I: PREREQUISITES

Objective

Handle the page fault exception by retrieving the
faulted page **from H.D.D to memory**



Load Program [env_create]



THREE kernel dynamic allocations:

1. **`create_page_table()`:** create new page table and link it to directory.
2. **`create_user_directory()`:** create new user directory.
3. **`create_user_kern_stack(...)`:** create new user kernel stack.

Refer to APPENDICES for:

- Page File Helper Functions
- Working Set Structure & Helper Functions

Working Set: Structure

inc/environment_definitions.h

```
struct Env {  
    //...  
    //=====/*WORKING SET*/=====//page working set management  
    struct WS_List page_WS_list;    FIFO, CLK, ModCLK //List of WS elements  
    struct WorkingSetElement* page_last_WS_element; //ptr to last inserted WS element  
    unsigned int page_WS_max_size; //Max allowed size of WS
```

Proc Limit

- Each process has a **working set LIST** that is initialized in **env_create()**
- Its **max size** is set in "**page_WS_max_size**" during the **env_create()**
 - "page_last_WS_element" will point to either:
 - the **next location** in the WS after the last set one If **list is full**.
 - Null if the list is **not full**.

Each Element

inc/environment_definitions.h

```
struct WorkingSetElement {  
    unsigned int virtual_address;  
    unsigned int time_stamp; LRU  
    unsigned int sweeps_counter;  
    LIST_ENTRY(WorkingSetElement) prev_next_info;
```

- This list hold pointers to **struct** containing info about the currently loaded pages in memory.
- Each struct holds two important values about each page:
 1. User virtual address of the page
 2. Previous & Next pointers to be used by list

Working Set: Functions [GIVEN]

```
struct WorkingSetElement* env_page_ws_list_create_element  
    (struct Env* e, uint32 virtual_address)
```

Description:

1. **Allocate** a new object of `struct WorkingSetElement` (using `kmalloc()`)
2. **Initialize** it by the given virtual address

Return

1. On success: pointer to the created object
2. On failure: kernel should `panic()`

Working Set: Functions [GIVEN]

```
void env_page_ws_print(struct Env* e)
```

Description:

- Print the page working set **virtual addresses** together with **used**, **modified** & **buffered** bits.
- It also shows where the **page_last_WS_element** of the working set is point to

Working Set: Functions [GIVEN]

```
void env_page_ws_invalidate(struct Env* e, uint32  
                           virtual_address)
```

Description:

- Search for the given virtual address inside the working set of “e”, if found:
 1. Remove its WS element from the **list**
 2. Delete this element from the kernel **memory** (using **kfree()**)
 3. Unmap the page from the address space

#9: Kernel Dyn. Alloc. for a Process

```
void* create_user_kern_stack(uint32* ptr_user_page_directory)
```

Description:

1. **Create** a user kernel stack of size *KERNEL_STACK_SIZE*
2. **Mark** its **bottom** page as **NOT PRESENT** (GUARD page)

Return

1. On success: pointer to the created stack
2. On failure: kernel should **panic()**

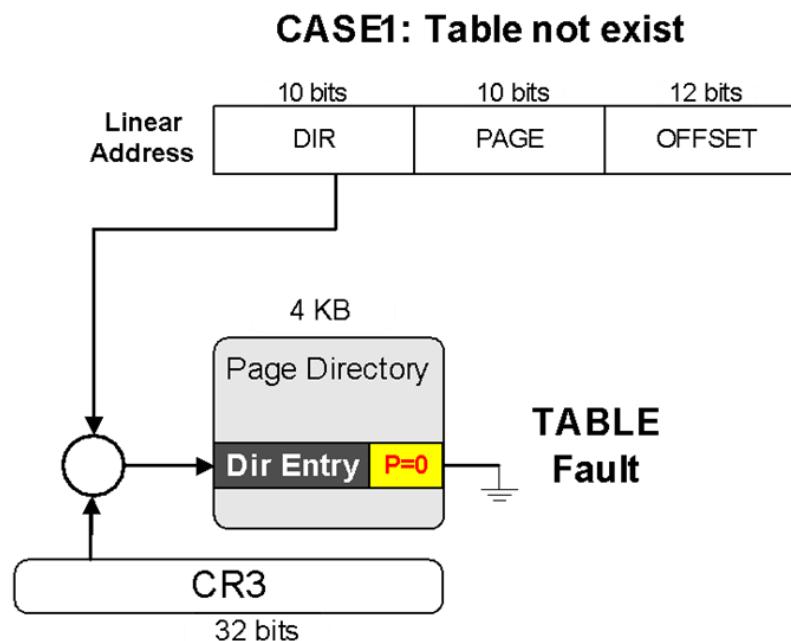
Fault Handler I: Overview

The main functions required to handle “**Page Fault**” are:

#	Function	File
1	<code>fault_handler</code>	Functions definitions <u>TO DO</u> in: kern/trap/fault_handler.c
2	<code>page_fault_handler</code>	

Fault Handler I: Overview

- **Fault:** is an exception thrown by the processor (MMU) to indicate that:
 - A **page table** not exist in the main memory (i.e. new table). (see the following figure) **OR**
 - A **page** can't be accessed due to either it's not present in the main memory



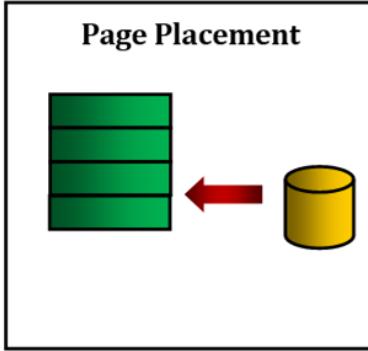
#10: Check Invalid Pointers

```
void fault_handler(struct Trapframe *tf)
```

Description:

- Check the **faulted_va** on the following:
 1. Pointing to **UNMARKED** page in user heap (i.e. its **PERM_UHPAGE** should equals 0)
 2. Pointing to **kernel**
 3. Exist with **read-only** permissions
- If **invalid** (any of above occur): it must be rejected without harm to the kernel or other running processes, by **exiting** the process using **env_exit()**

#11: Placement



Refer to APPENDICES for:



Page File Helper Functions



Working Set Struct & Helper Func's

`page_fault_handler(struct Env * faulted_env, uint32 fault_va)`

If the size of the page working LIST < *its max size*, then do (refer to appendices for helper functions)

Scenario 1: Placement

1. **Allocate** space for the faulted page
2. **Read** the faulted page from page file to memory
3. If the page **does not exist** on page file, then
 1. If it is a **stack** or a **heap** page, then, it's OK.
 2. Else, it must be **rejected** without harm to the kernel or other running processes, by **exiting** the process.
4. Reflect the changes in the page working set list (i.e. add new element to list & update its last one)

NEXT, IT'S YOUR TURN...

😊 Enjoy **developing** your **own OS** 😊

